



中国软件技术大会

ByteIR: 迈向端到端的AI编译

字节跳动 机器学习系统工程师 刘渊强

2023.12




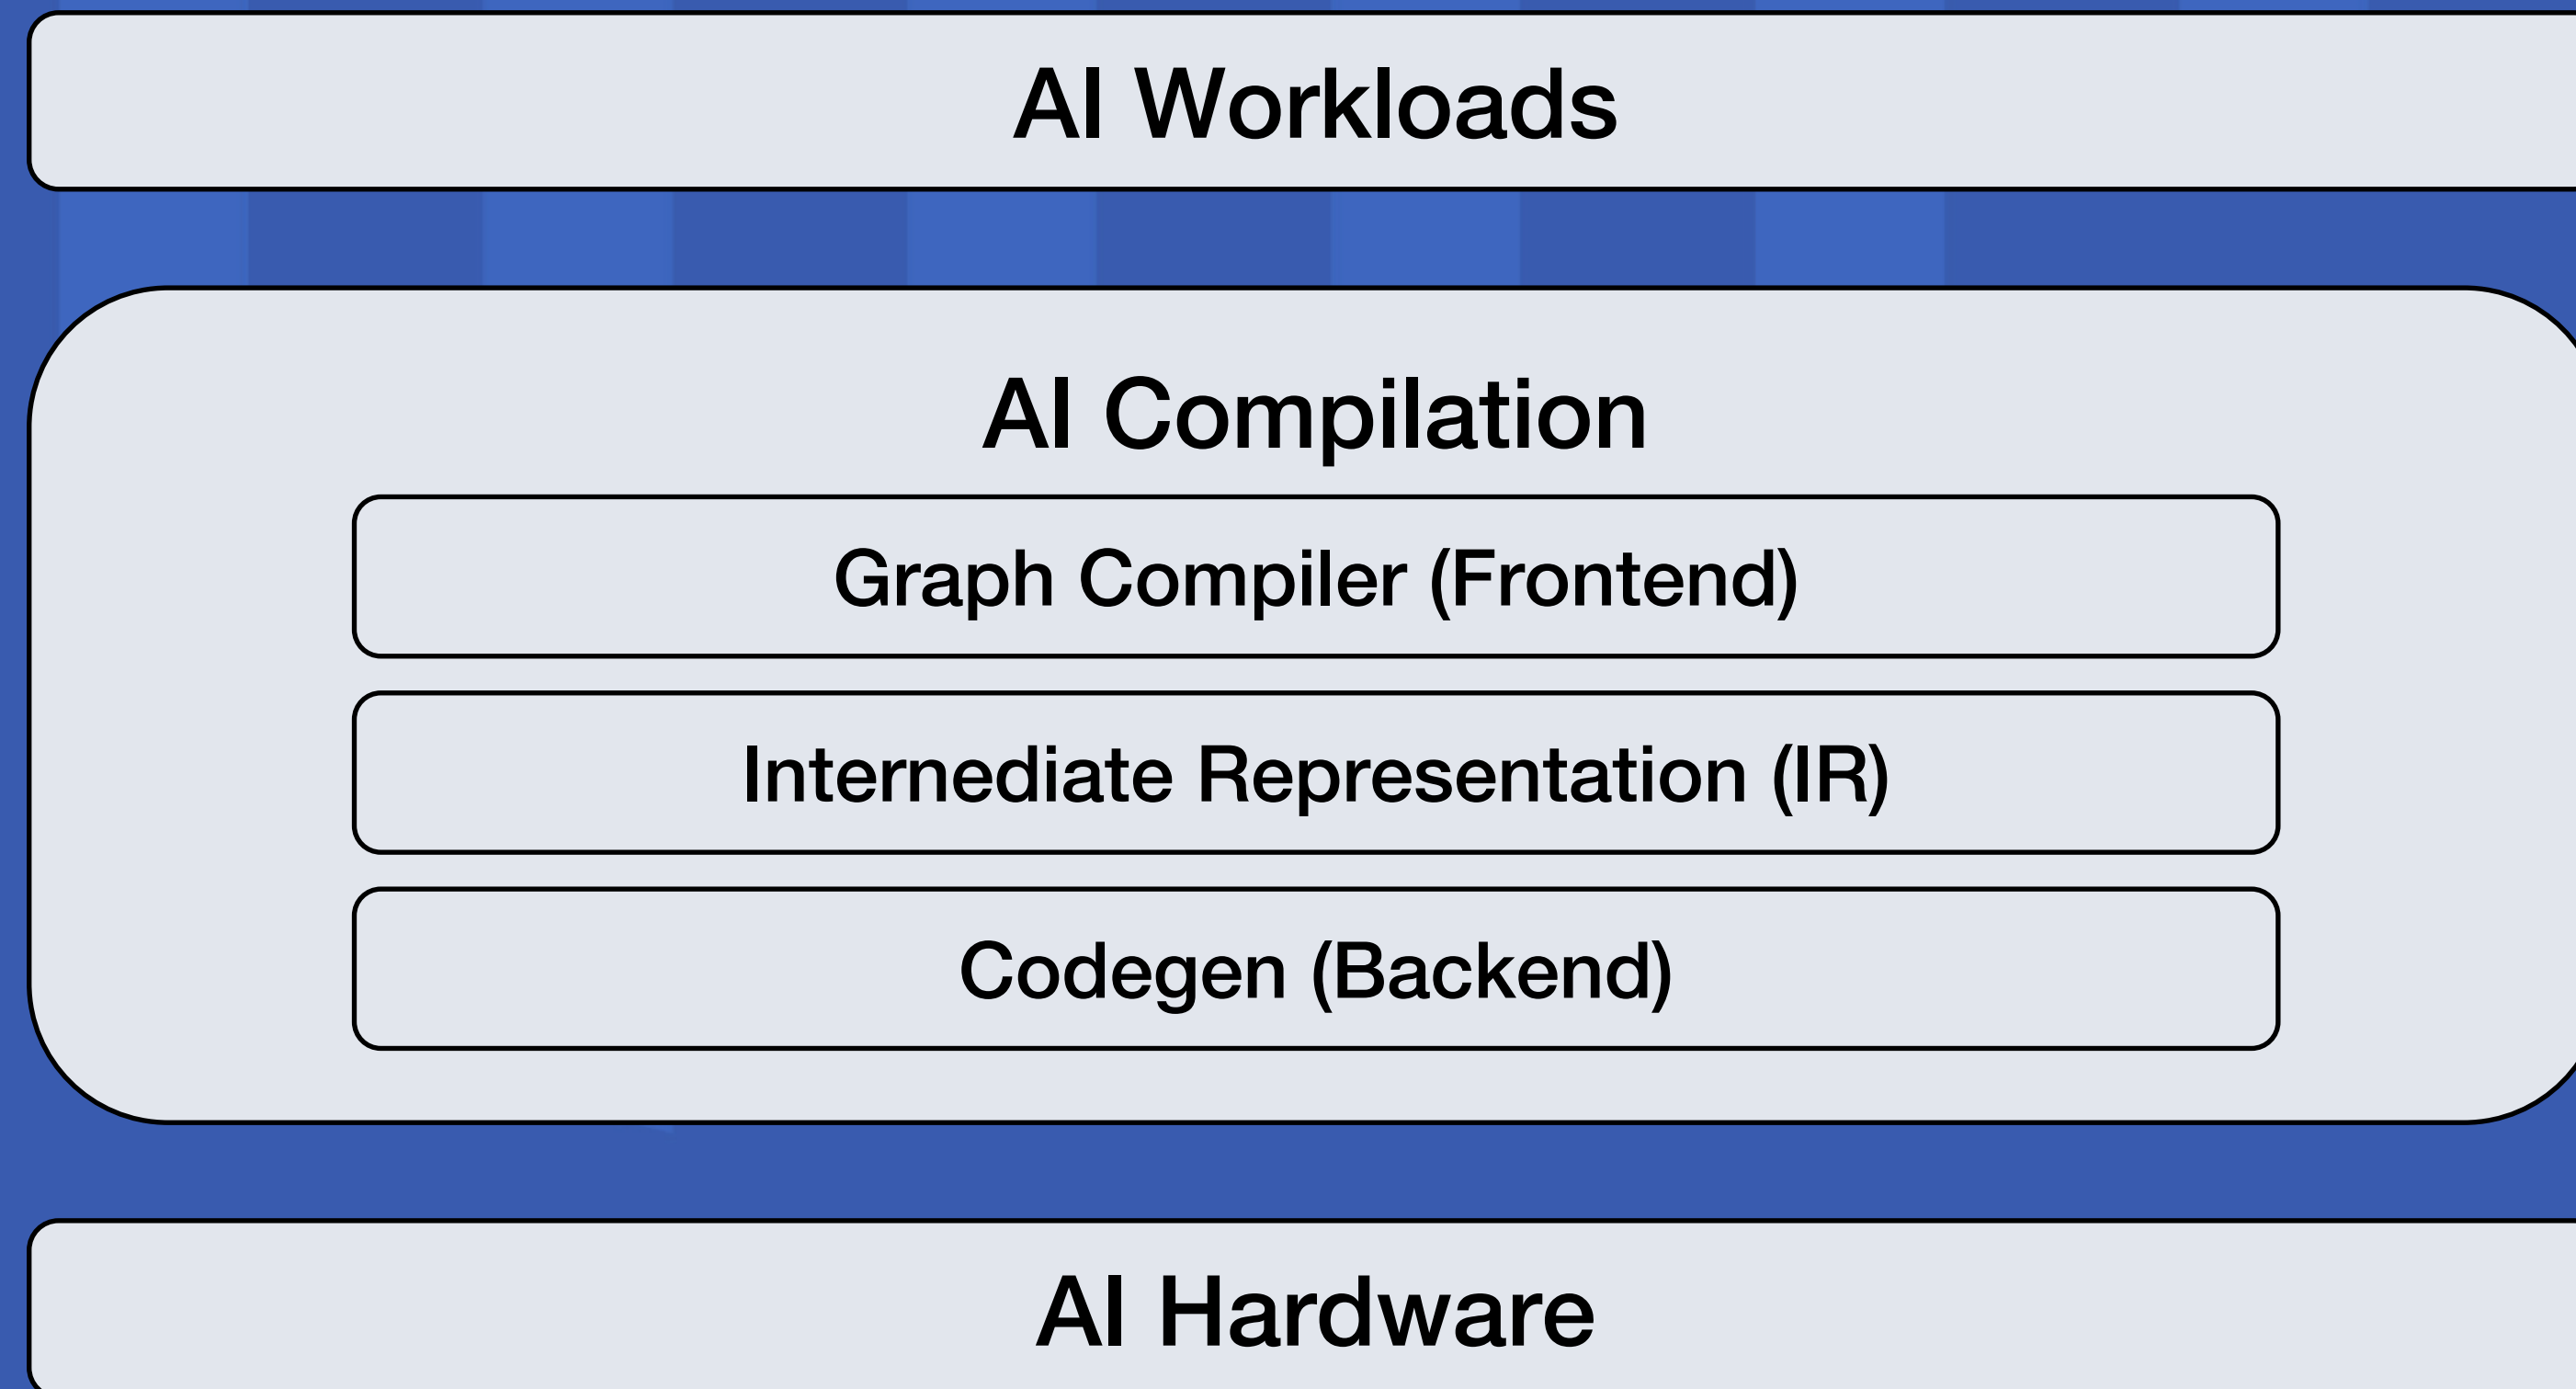
- 1. What is Byteler and advantages of Byteler**
- 2. Design and technical details**
- 3. LLM training example and performance**

What is ByteIR

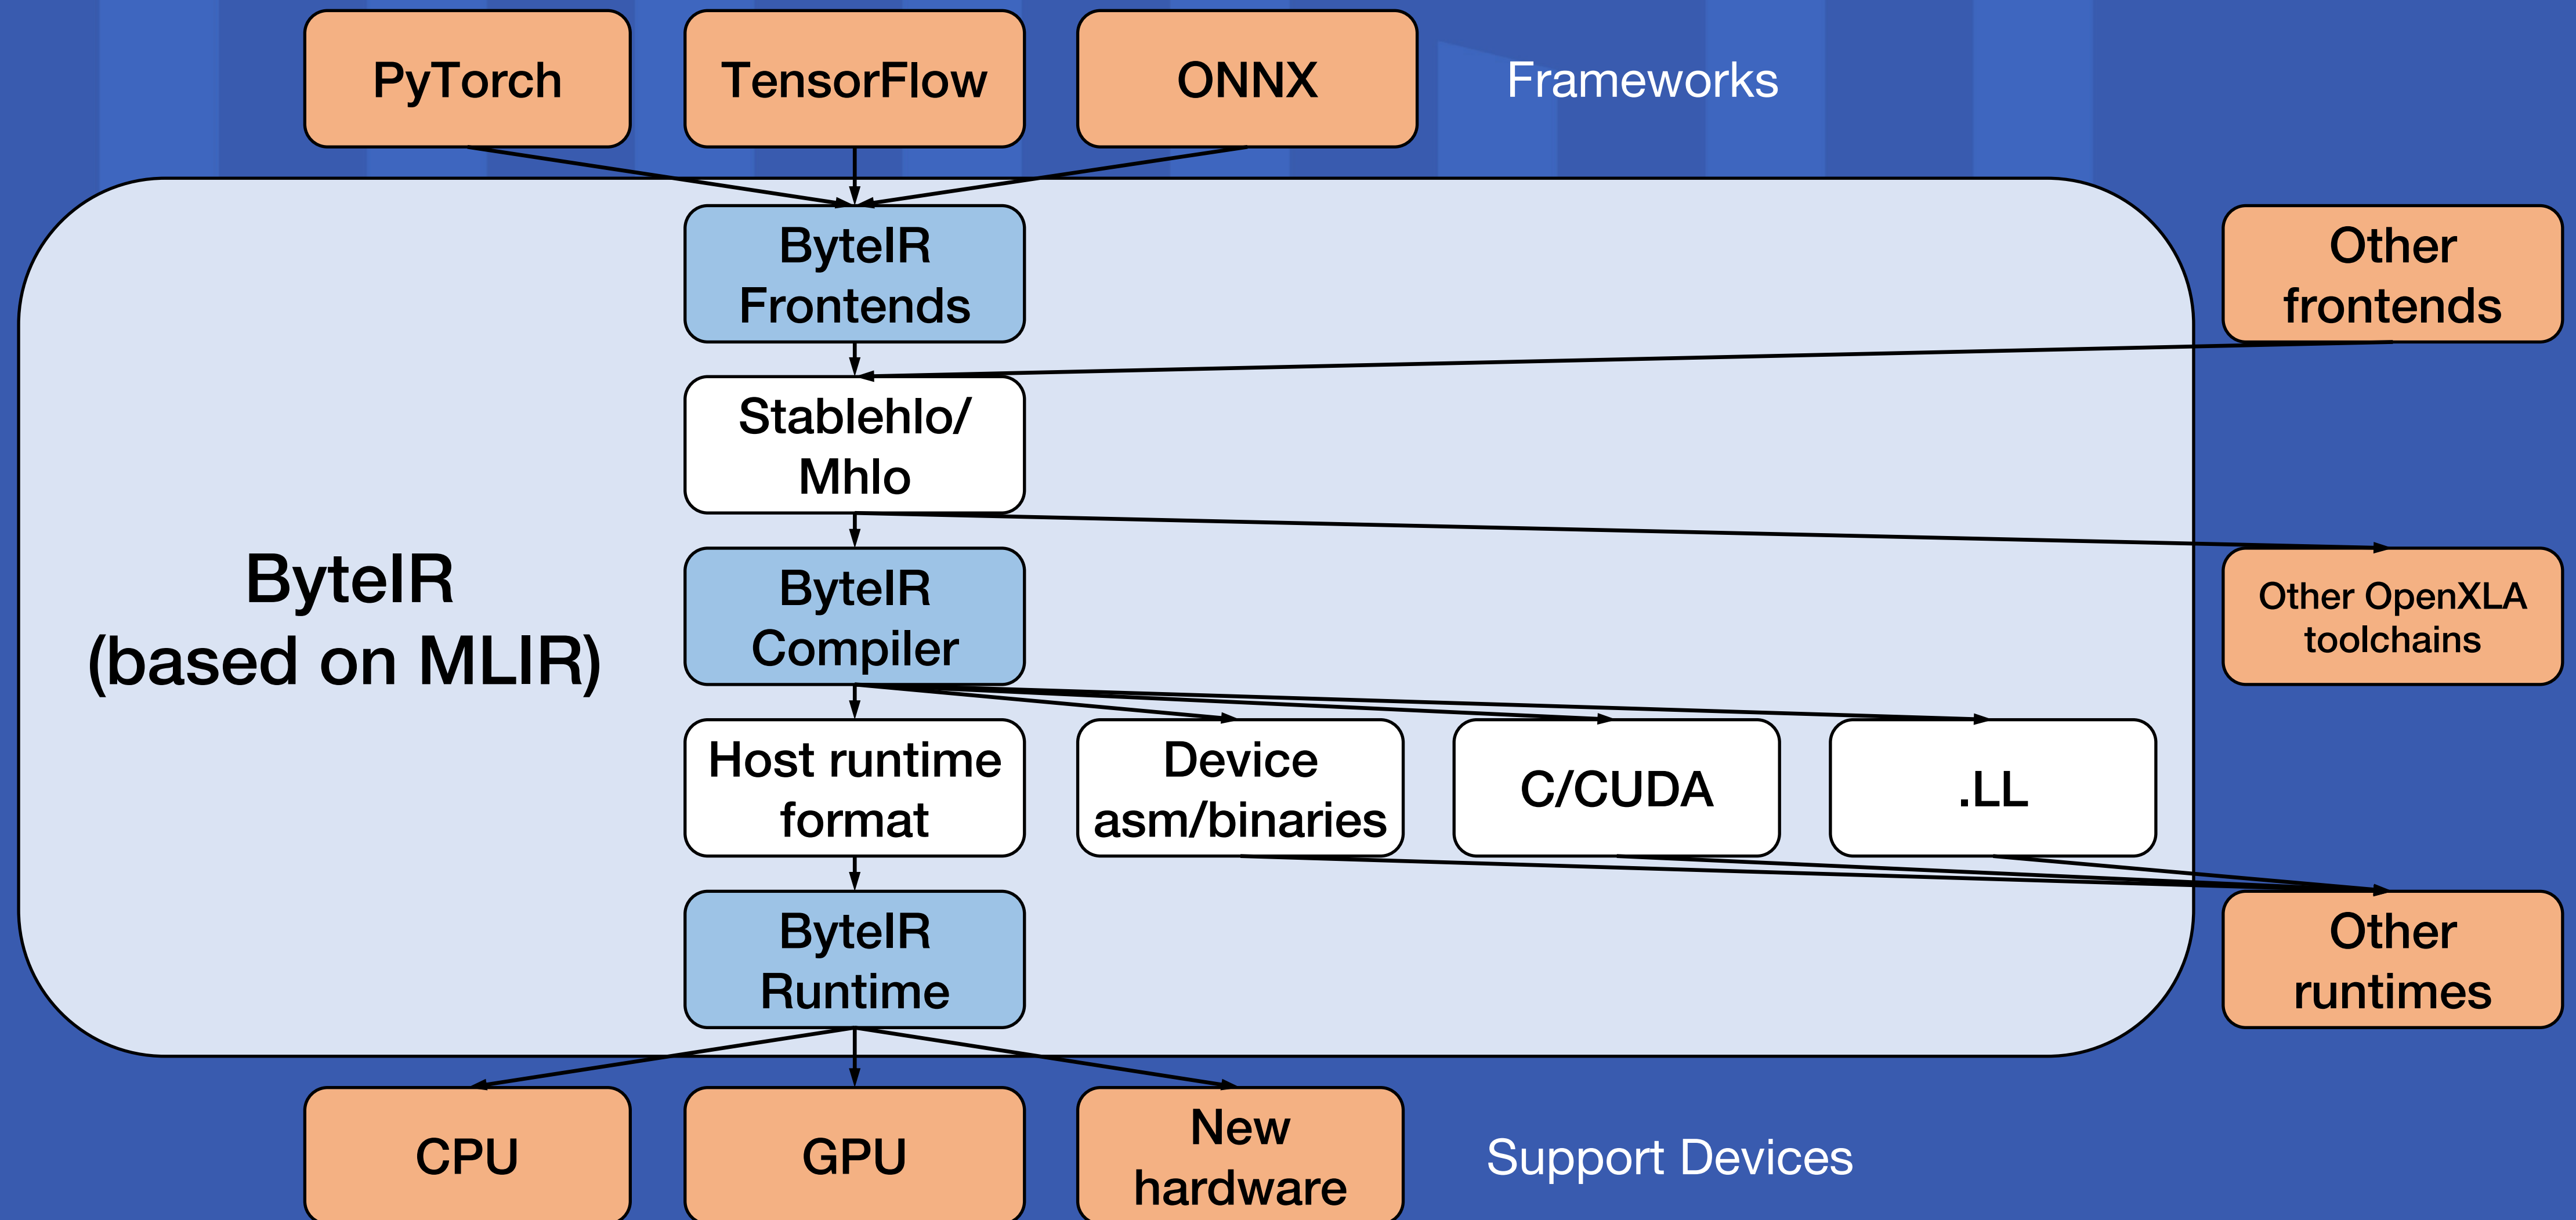
 BYTEIR is **NOT** a specific IR (Intermediate Representation)

 BYTEIR is our solution for framework-to-hardware compilation

AI Compilation: NN graph to HW



ByteIR Architecture



Advantages of BytelR

1. Embrace open source, upstream first

Contribute to llvm, tensorflow, pytorch, torch-mlir, onnx-mlir, stablehlo

2. Well support for PyTorch, both inference and training

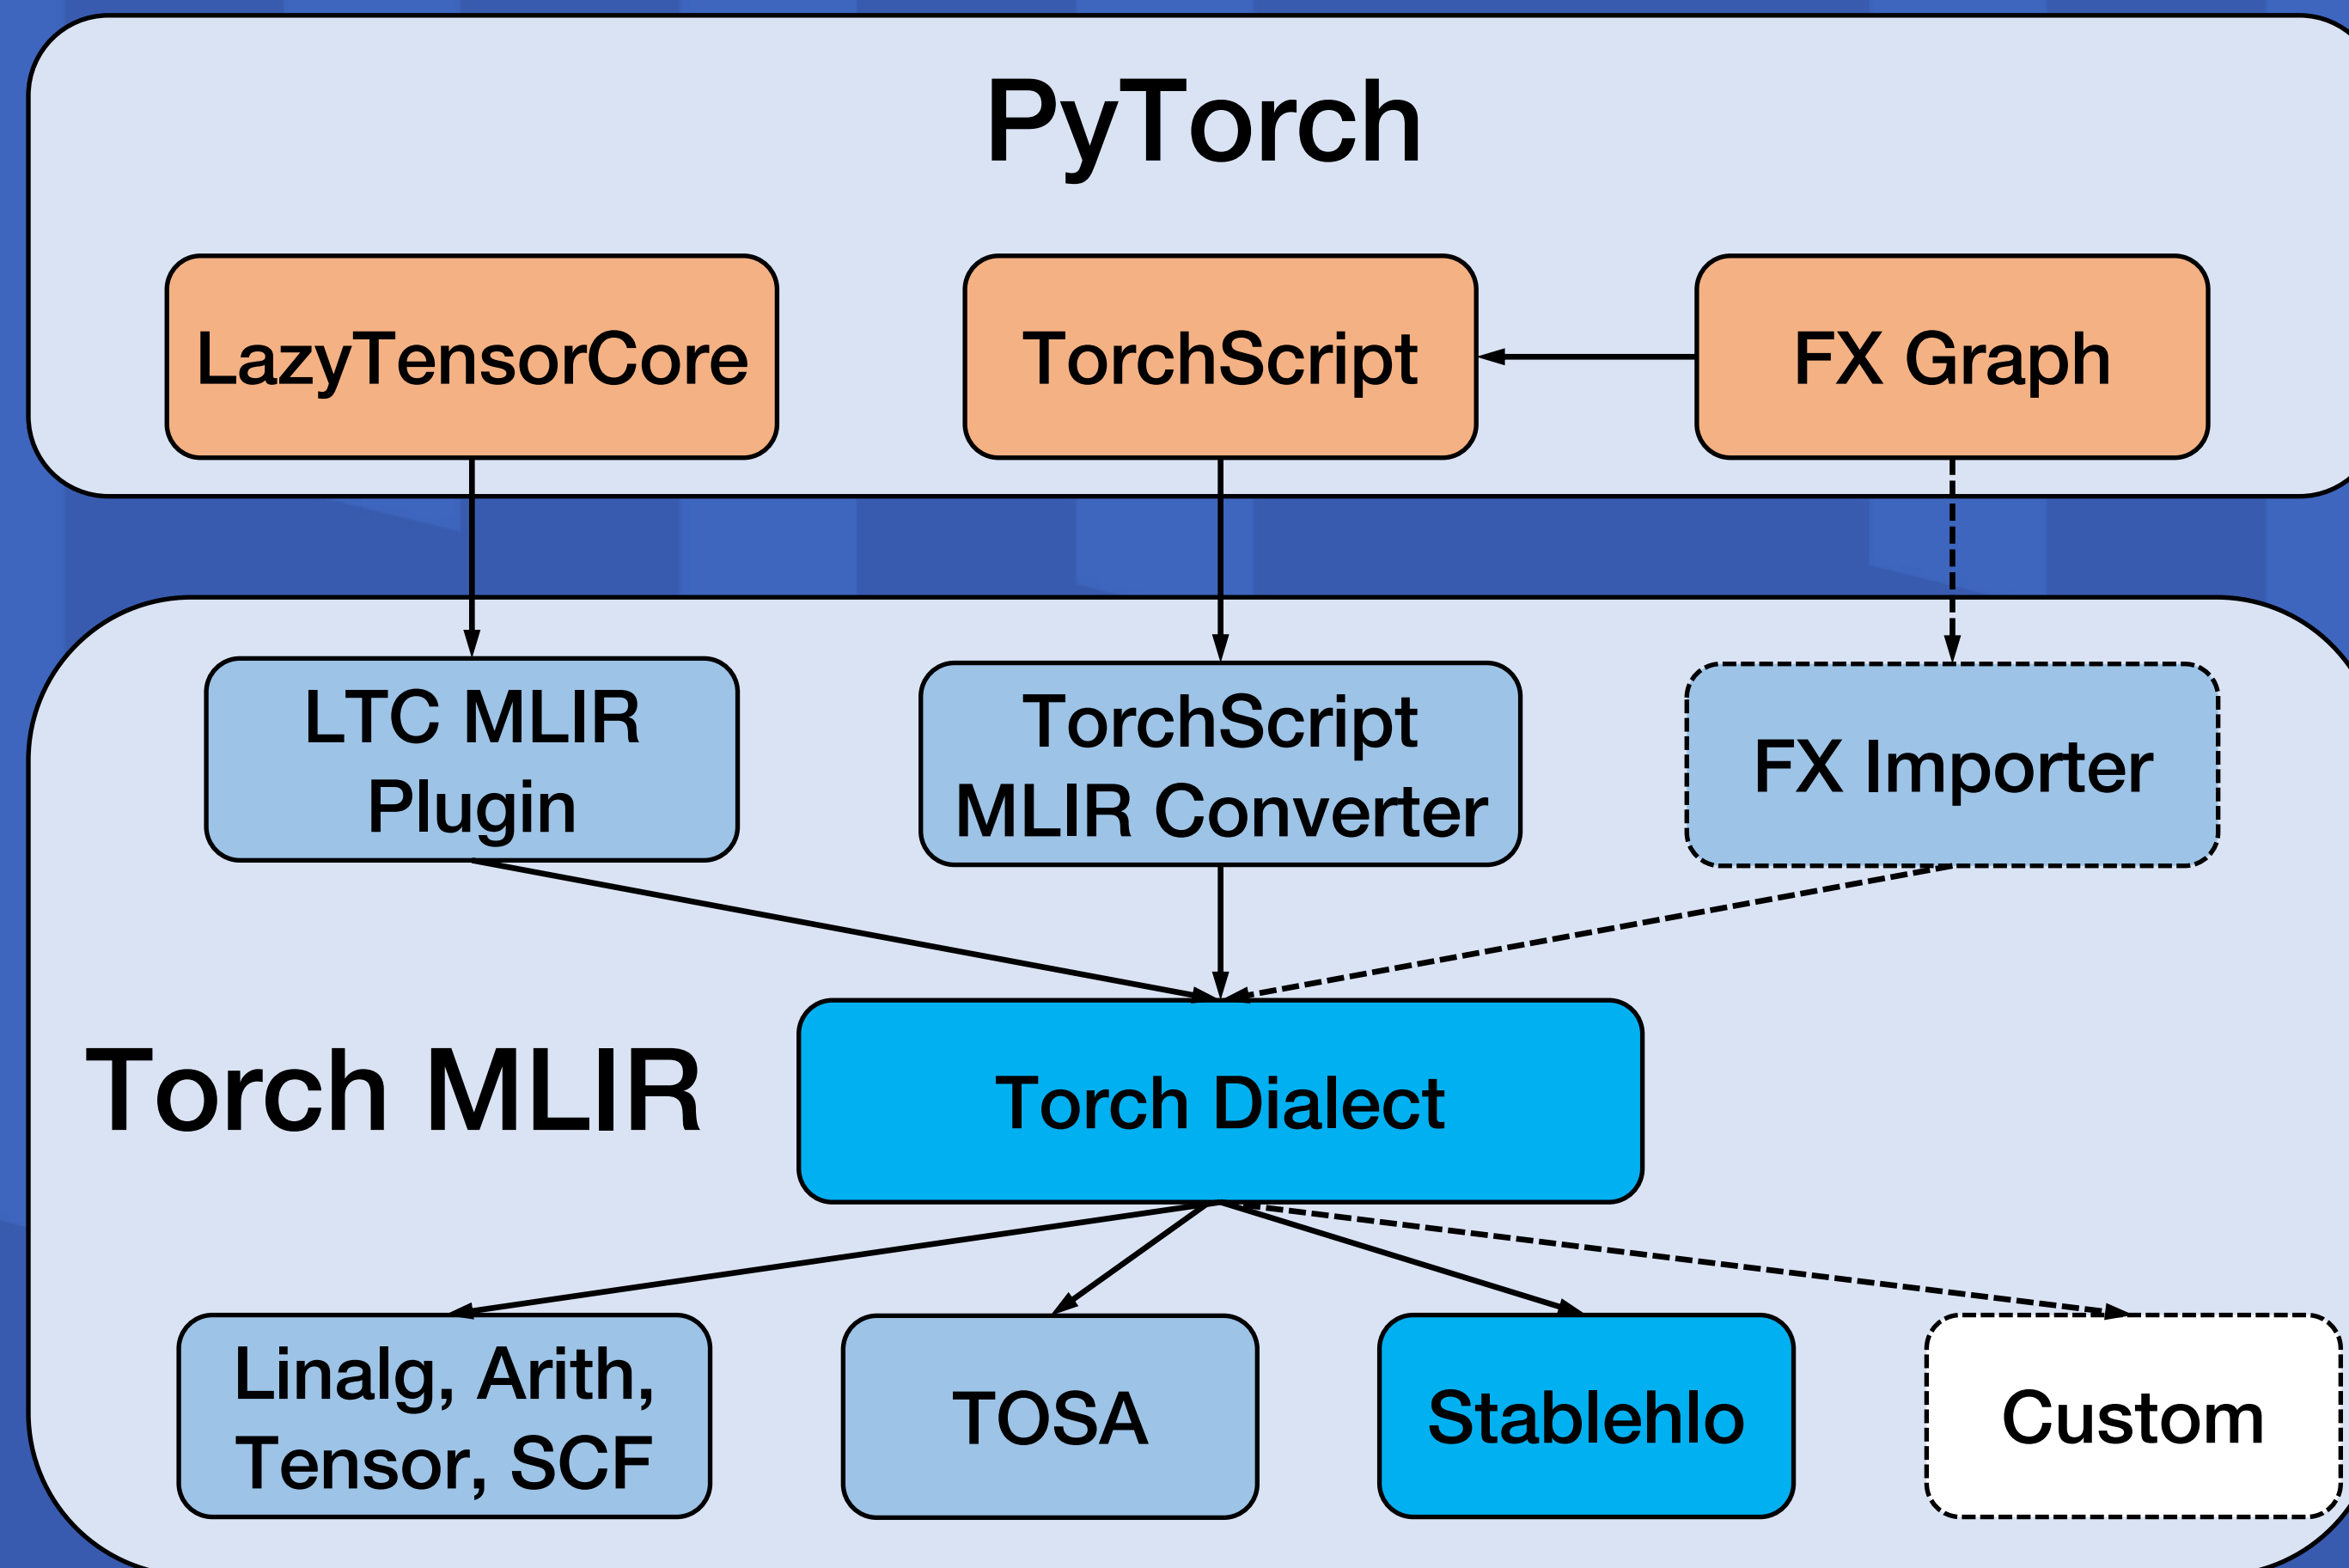
3. Friendly to new hardware (ASIC/NPU)

4. Flexible, extensible, high performance



1. What is BytelR and advantages of BytelR
2. Design and technical details
3. LLM training example and performance

ByteIR Torch Frontend – Torch MLIR



Github: <https://github.com/llvm/torch-mlir>

Torch MLIR Lowering

torch

```
class Module(torch.nn.Module):  
    def __init__(self):  
        super().__init__()  
        self.linear = nn.Linear(10, 20)  
    def forward(self, x):  
        return self.linear(x)
```

torch dialect

```
func.func @forward(%arg0: !torch.tensor,  
%arg1: !torch.tensor) -> !torch.tensor {  
    %0 = torch.aten.mm %arg0,  
    %arg1 : !torch.tensor, !torch.tensor -> !torch.tensor  
    return %0 : !torch.tensor  
}
```

torch.jit.script
MLIR importer

Shape/dtype
inference
Value semantics
analysis


torch dialect

```
func.func @forward(%arg0: !torch.vtensor<[20,  
10], f32>, %arg1: !torch.vtensor<[10, 20], f32>) -  
> !torch.vtensor<[20, 20], f32> {  
    %0 = torch.aten.mm %arg0,  
    %arg1 : !torch.vtensor<[20, 10],  
    f32>, !torch.vtensor<[10, 20], f32> -  
> !torch.vtensor<[20, 20], f32>  
    return %0 : !torch.vtensor<[20, 20], f32>  
}
```

stablehlo dialect

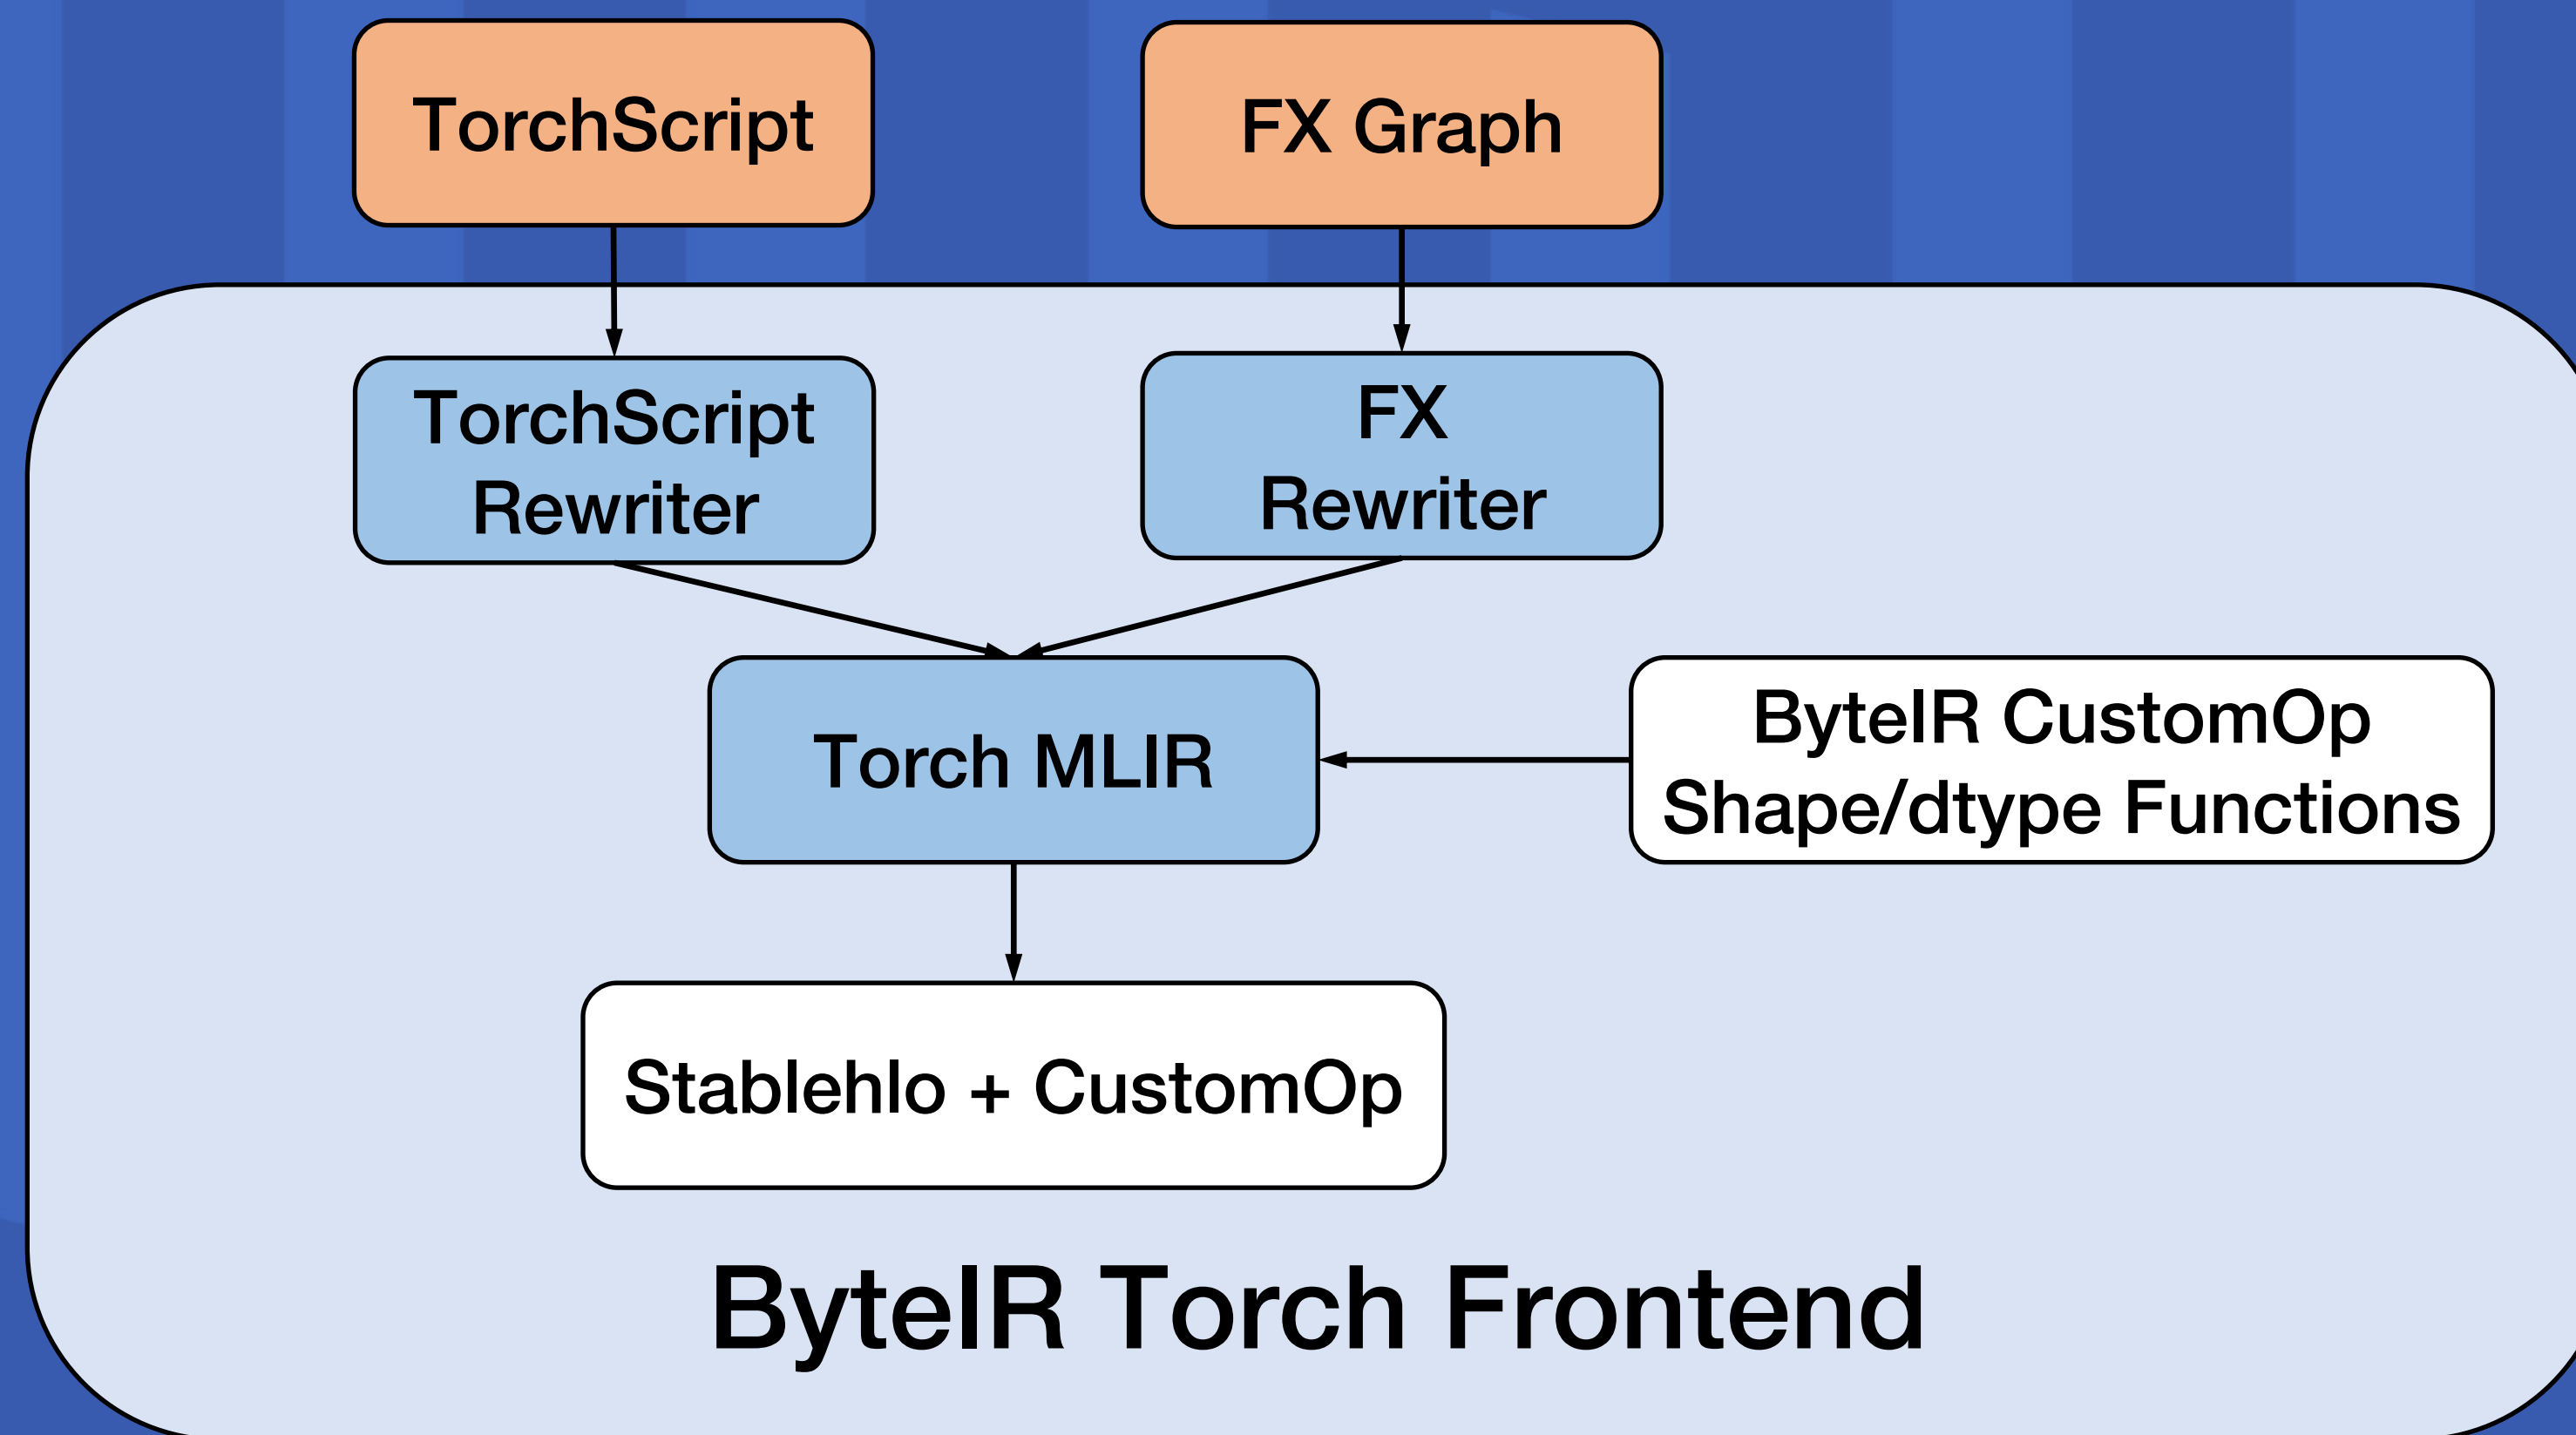
```
func.func @forward(%arg0: tensor<20x10xf32>,  
%arg1: tensor<10x20xf32>) ->  
tensor<20x20xf32> {  
    %0 = "stablehlo.dot"(%arg0, %arg1) :  
    tensor<20x10xf32>, tensor<10x20xf32> ->  
    tensor<20x20xf32>  
    return %0 : tensor<20x20xf32>  
}
```

Lowering



Decomposition

ByteIR Torch Frontend



Able to provide **coarse-grained** operators



Corner Cases

```
def forward(self, x):  
    x += 1  
    return x
```

Success, but no effect on input

```
def forward(self, x, y):  
    x += y  
    z = x.view(-1, 4)  
    x += 1  
    return x, z
```

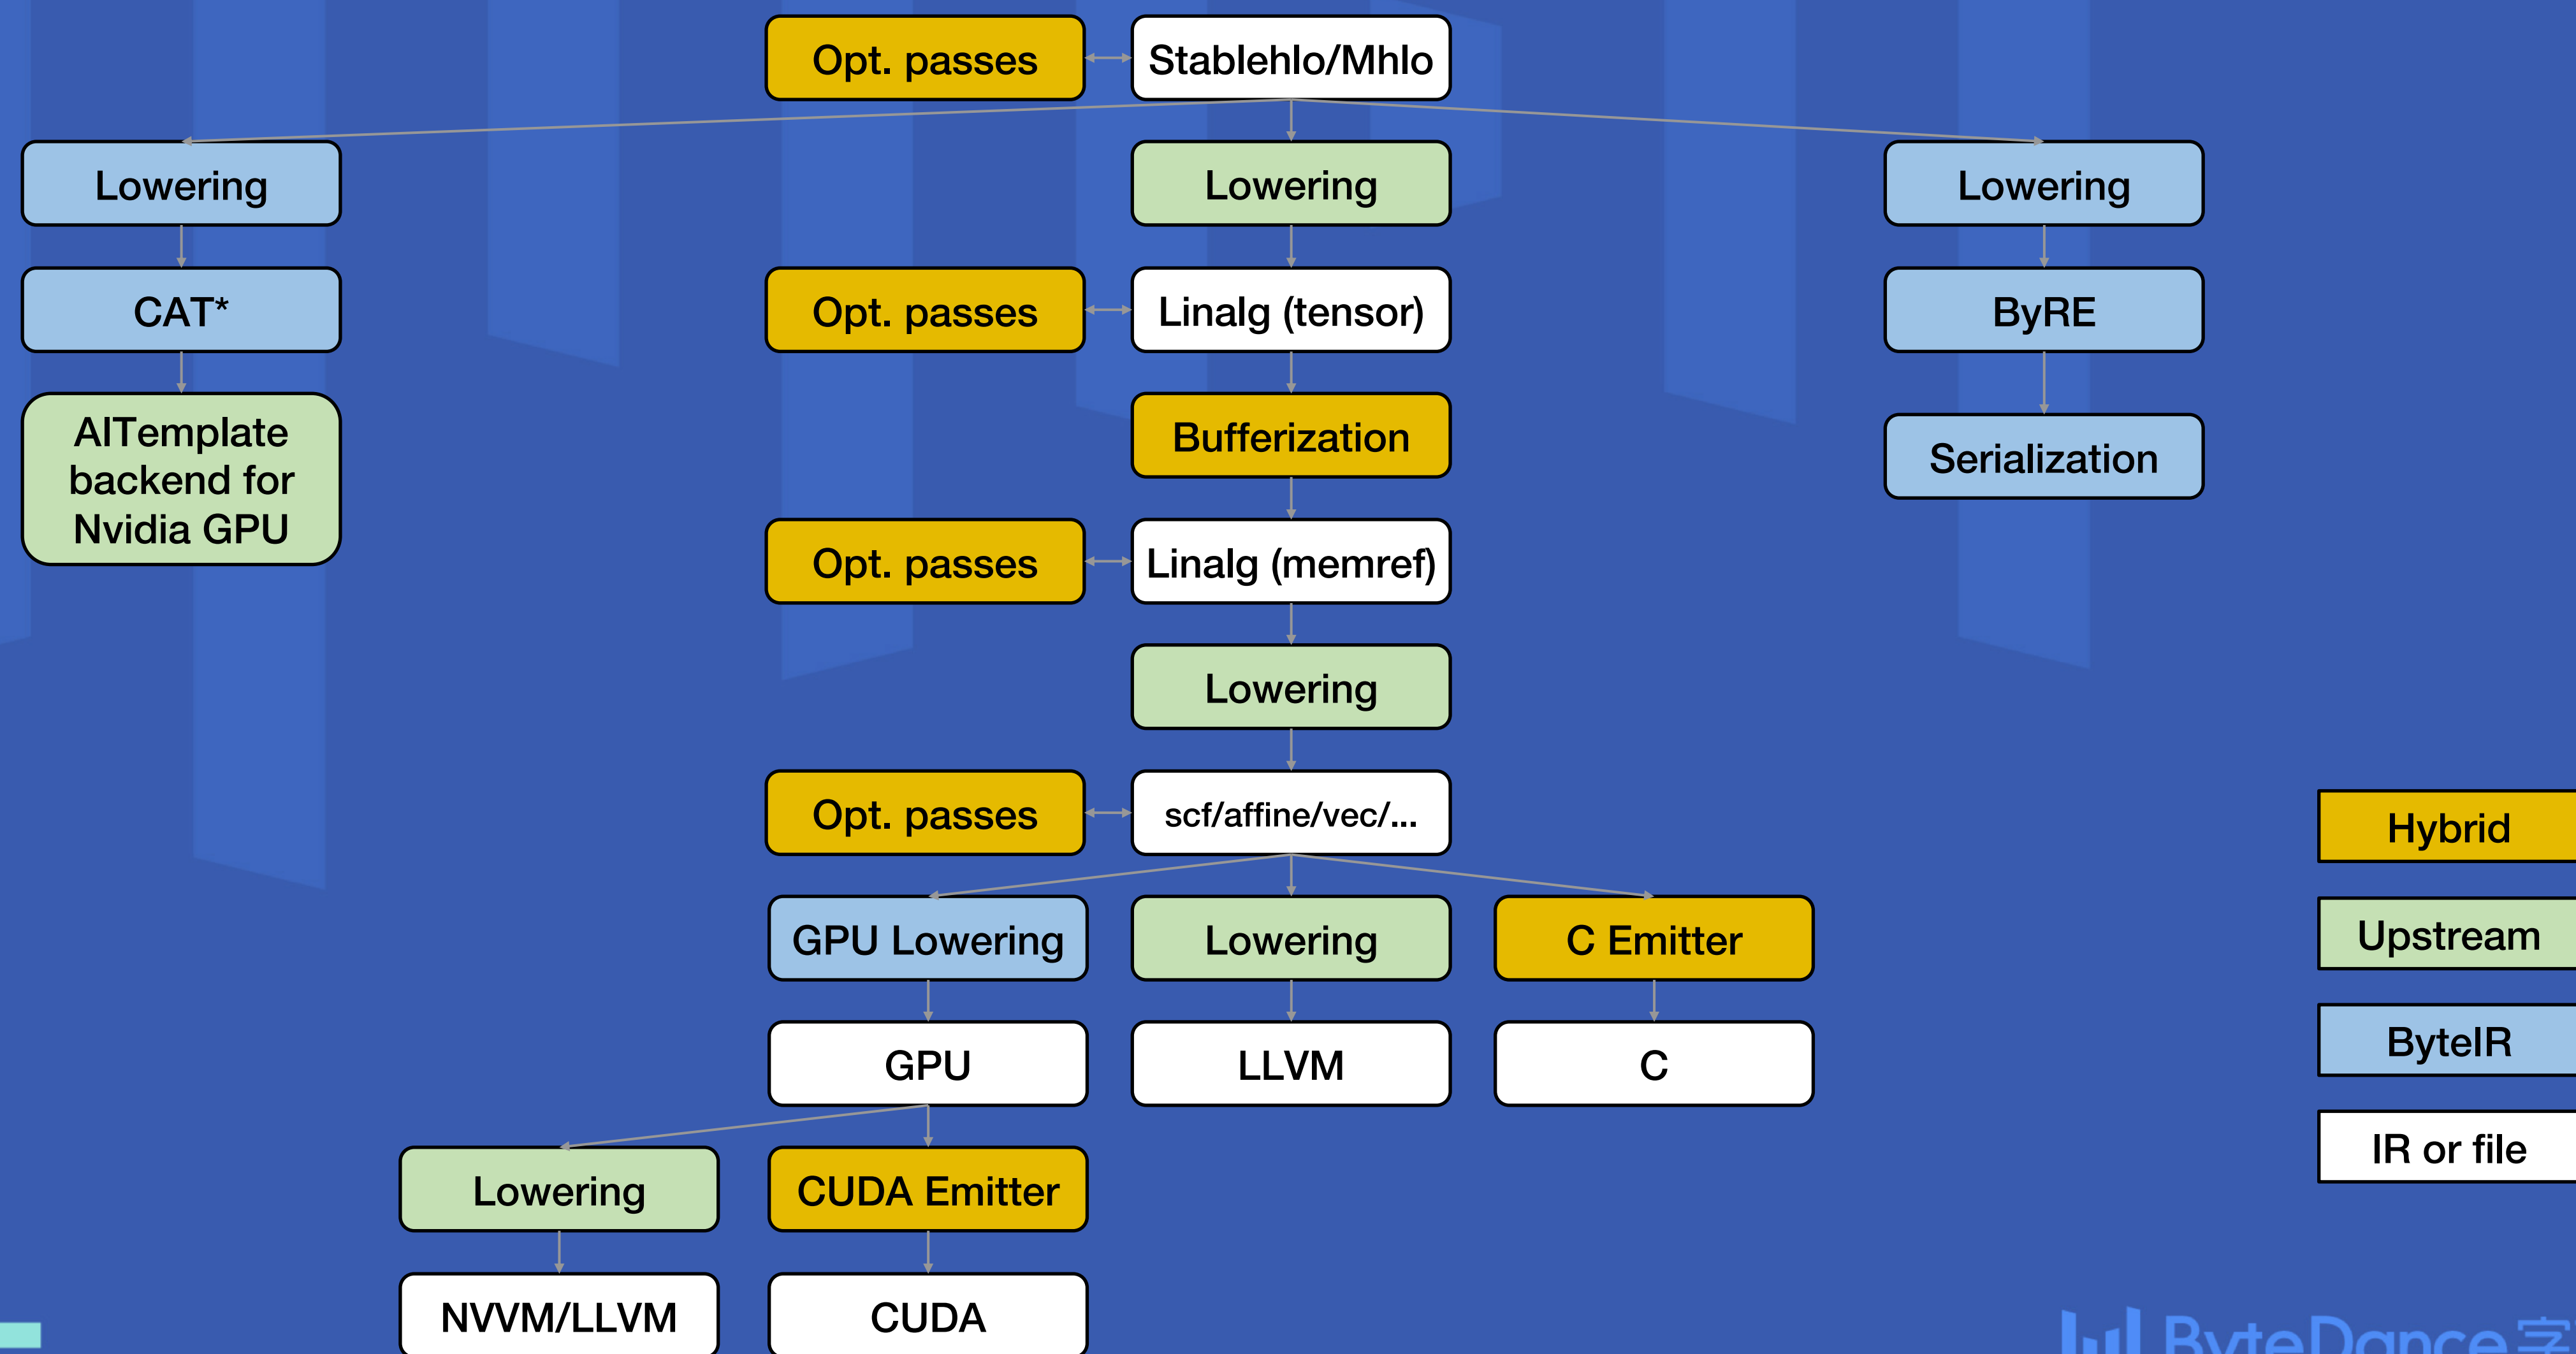
Failed, really overwrite same memory

```
def forward(self, x):  
    y = torch.as_strided(x, size, stride)  
    return y
```

Failed, no abstraction of pointer/storage

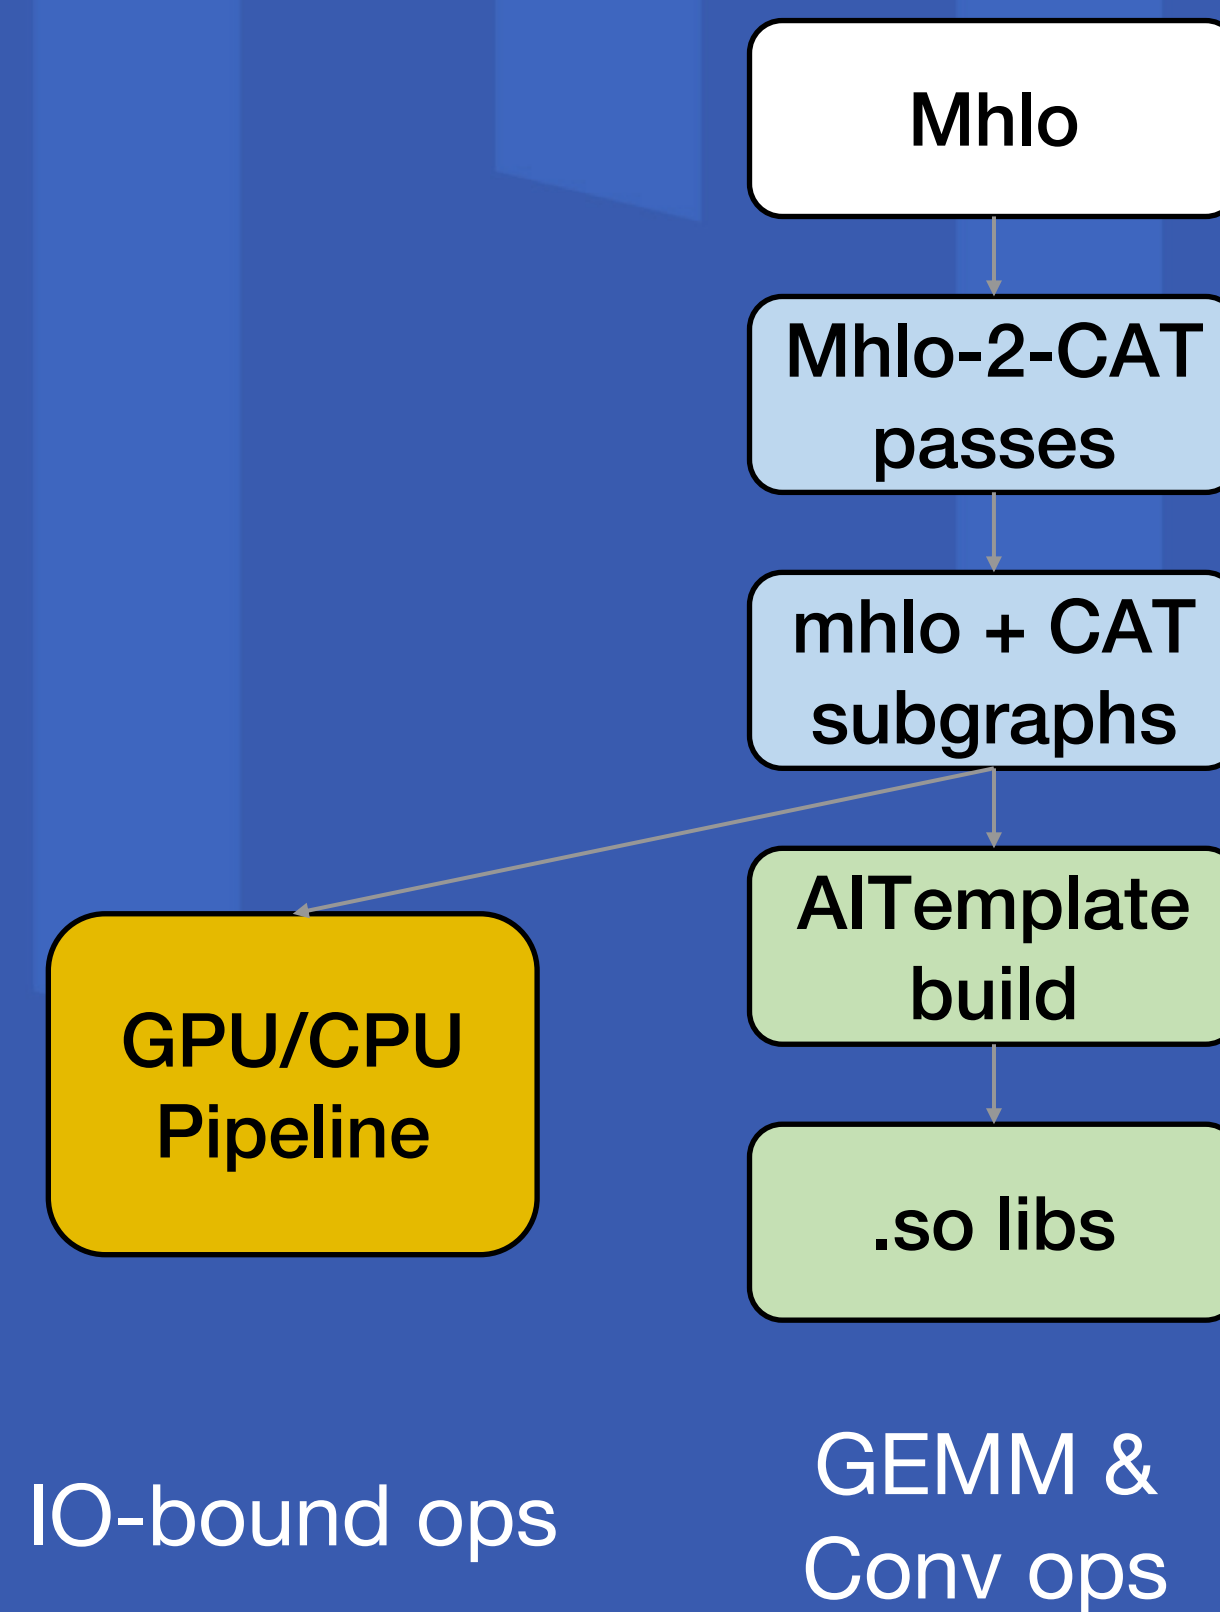
- Don't support dynamic if/for/while
.....

ByteIR Compiler Overview



Integrating AI Template

We introduce **CAT** (Composable Algebra Template) dialect



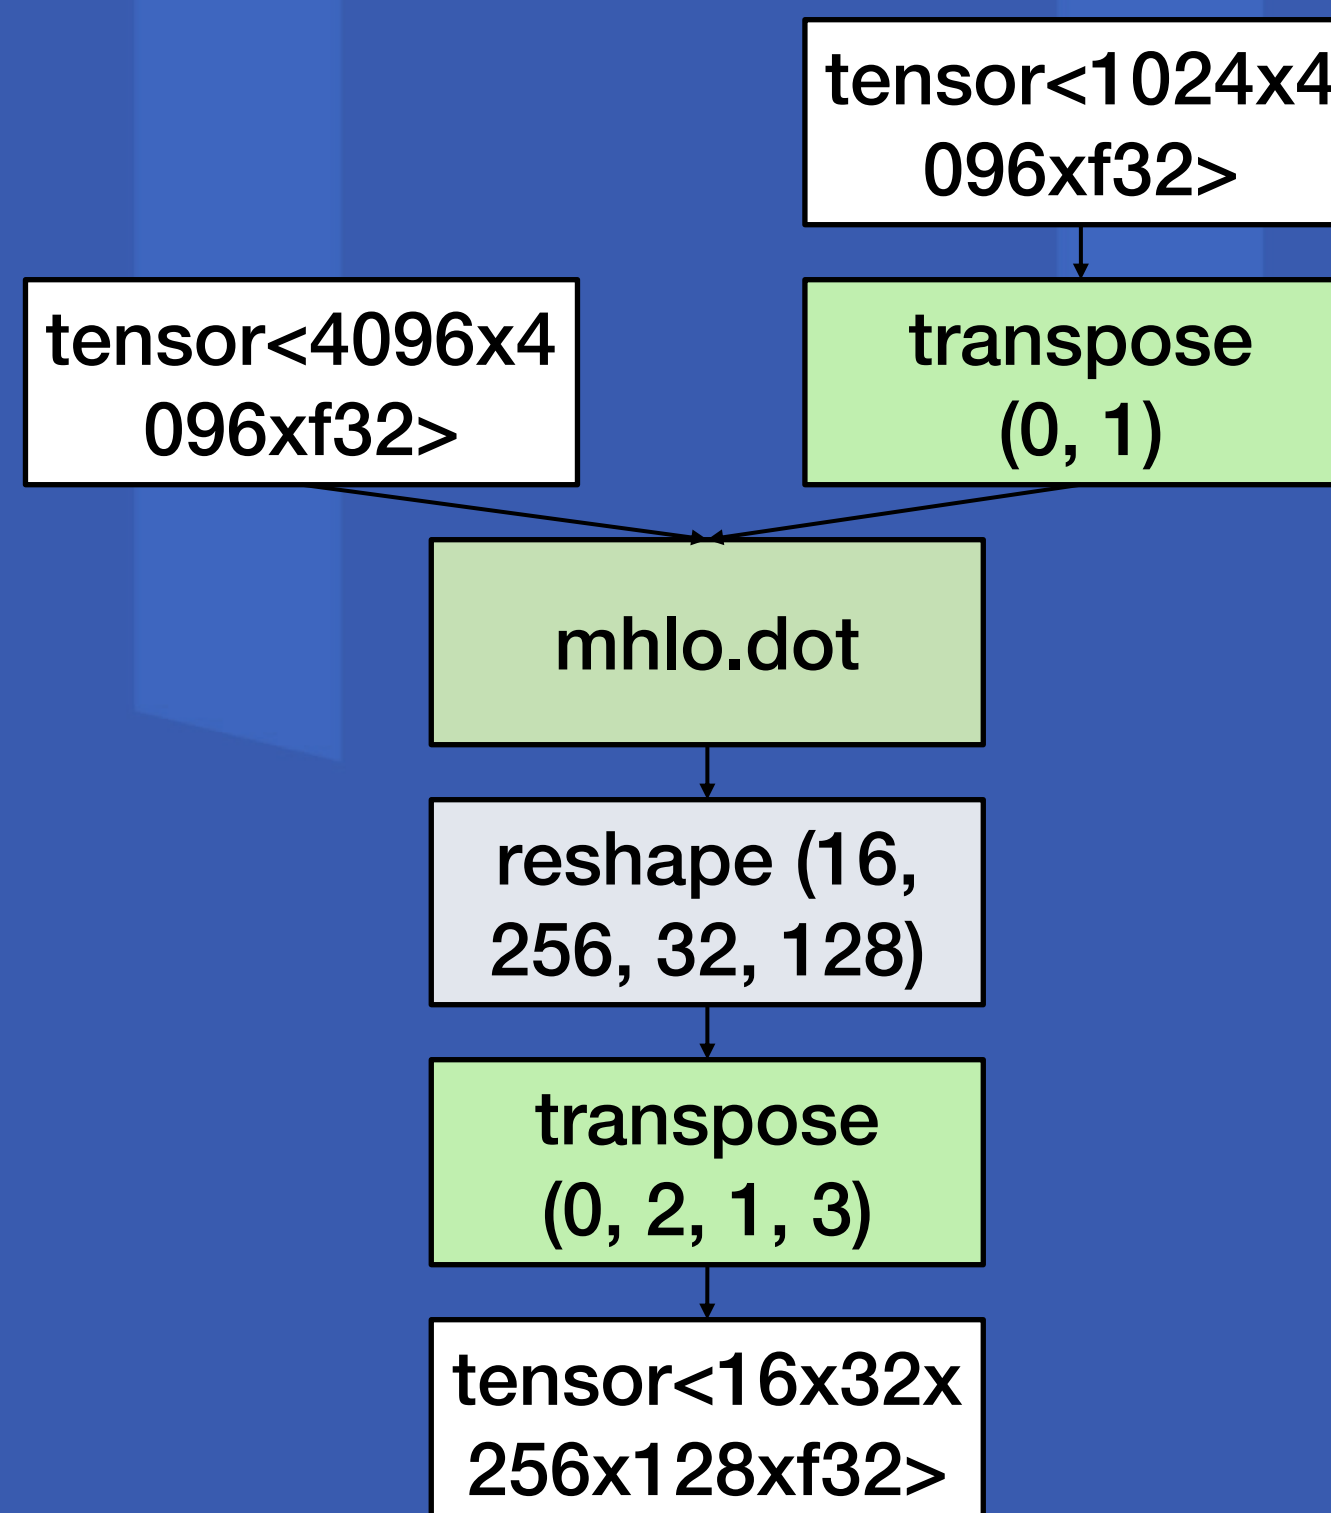
AI Template: Meta's inference framework converting PyTorch NN to CUDA code (based on CUTLASS)

Mhlo-2-CAT Passes


Mhlo-2-CAT:

- Convert Mhlo ops to CAT ops (one CAT op corresponds to one ALT op)
- Eliminate redundant transpose/permuate ops

Mhlo dialect



CAT dialect



Linalg Tiling and Fusion

Mhlo op:

```
func.func @fuse_element(%arg0 ..., %arg1 ...)
  %0 = mhlo.some_elemwise_binary_1(%arg0, %arg1)
  %1 = mhlo.some_elemwise_binary_2(%0, %arg1)
  return %1
```



Linalg transformation

Linalg op:

```
func.func @fuse_element(%arg0 ..., %arg1 ...)
  %1 = linalg.generic {indexing_maps = ...,
    iterator_types = ...} ins(%arg0, %arg1) outs... {
    ^bb0(%in0, %in1, %out)
      %2 = arith.some_elemwise_binary_1(%in0, %in1)
      %3 = arith.some_elemwise_binary_2(%2, %in1)
      linalg.yield %3 ...
  }
  return %1
```

ByeIR's Linalg Extension

More ops:

- Alias, Diag, Scan, Scatter, Softmax, TopK
- support transformations of extended ops

Enhanced fusion transformations:

- producer-consumer & input-sharing fusion
- tiling along reduction axis correction
- intermediates as outputs within fusion
- intermediate tensor dim simplification
- map ops to generic ops conversion
- ...

Other introduced transformations:

- Collapse dims transformation
- Fuse operands transformation
- ...

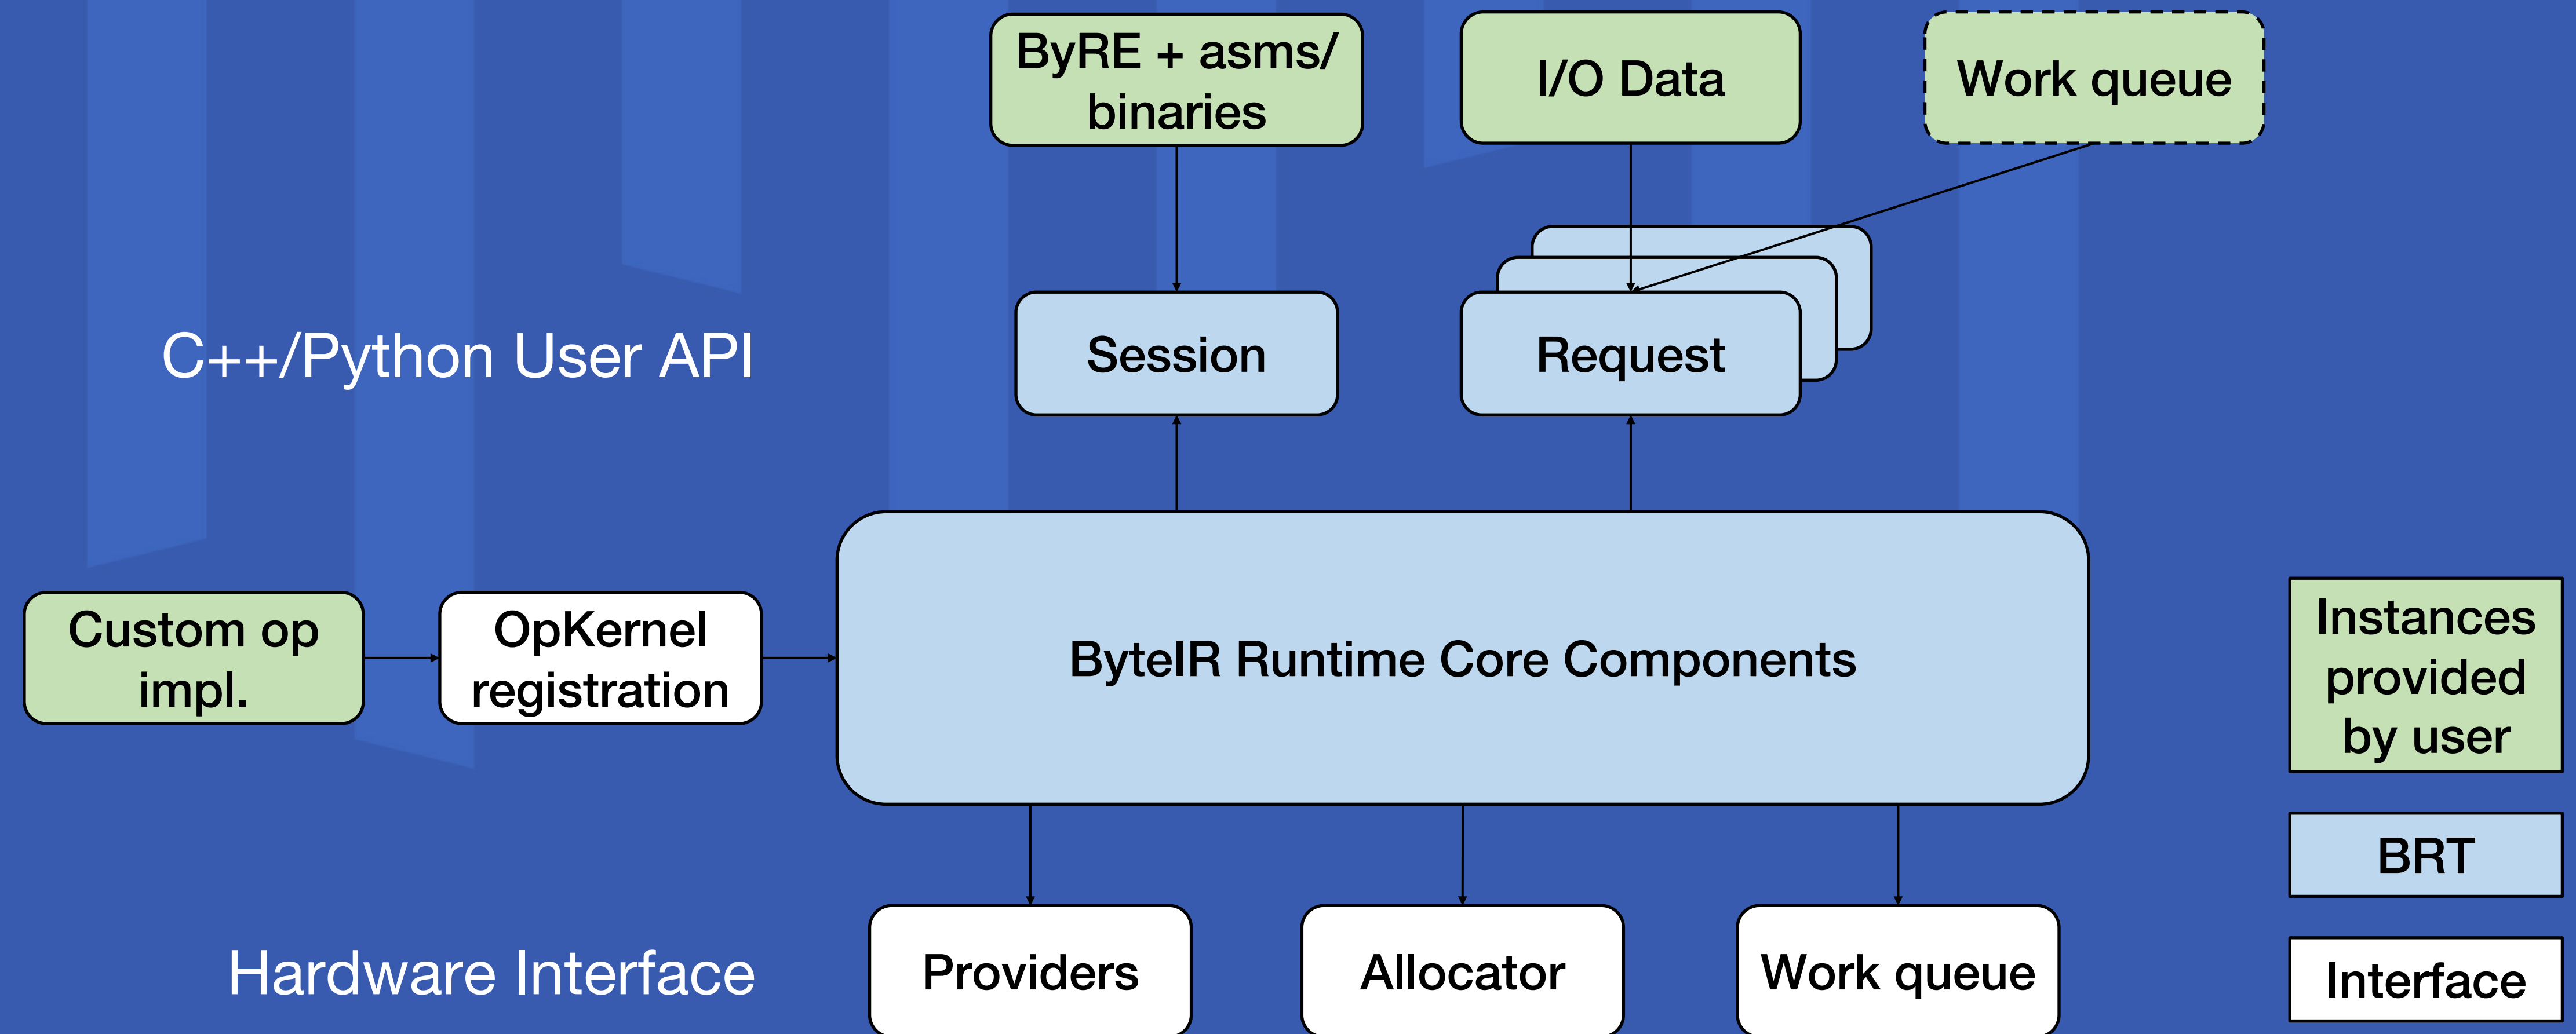


Benefits:

- Extreme IO-bound op fusion
- Lower overhead for fused ops (Exploit GPU DRAM bandwidth)



ByteIR Runtime (BRT) Overview



BRT Interface for Hardware

Provider

A collection of op implementation

- e.g., mm, d2h/h2d memcpy

Work Queue

Abstraction(like CUDAStream) for execution order

Allocator

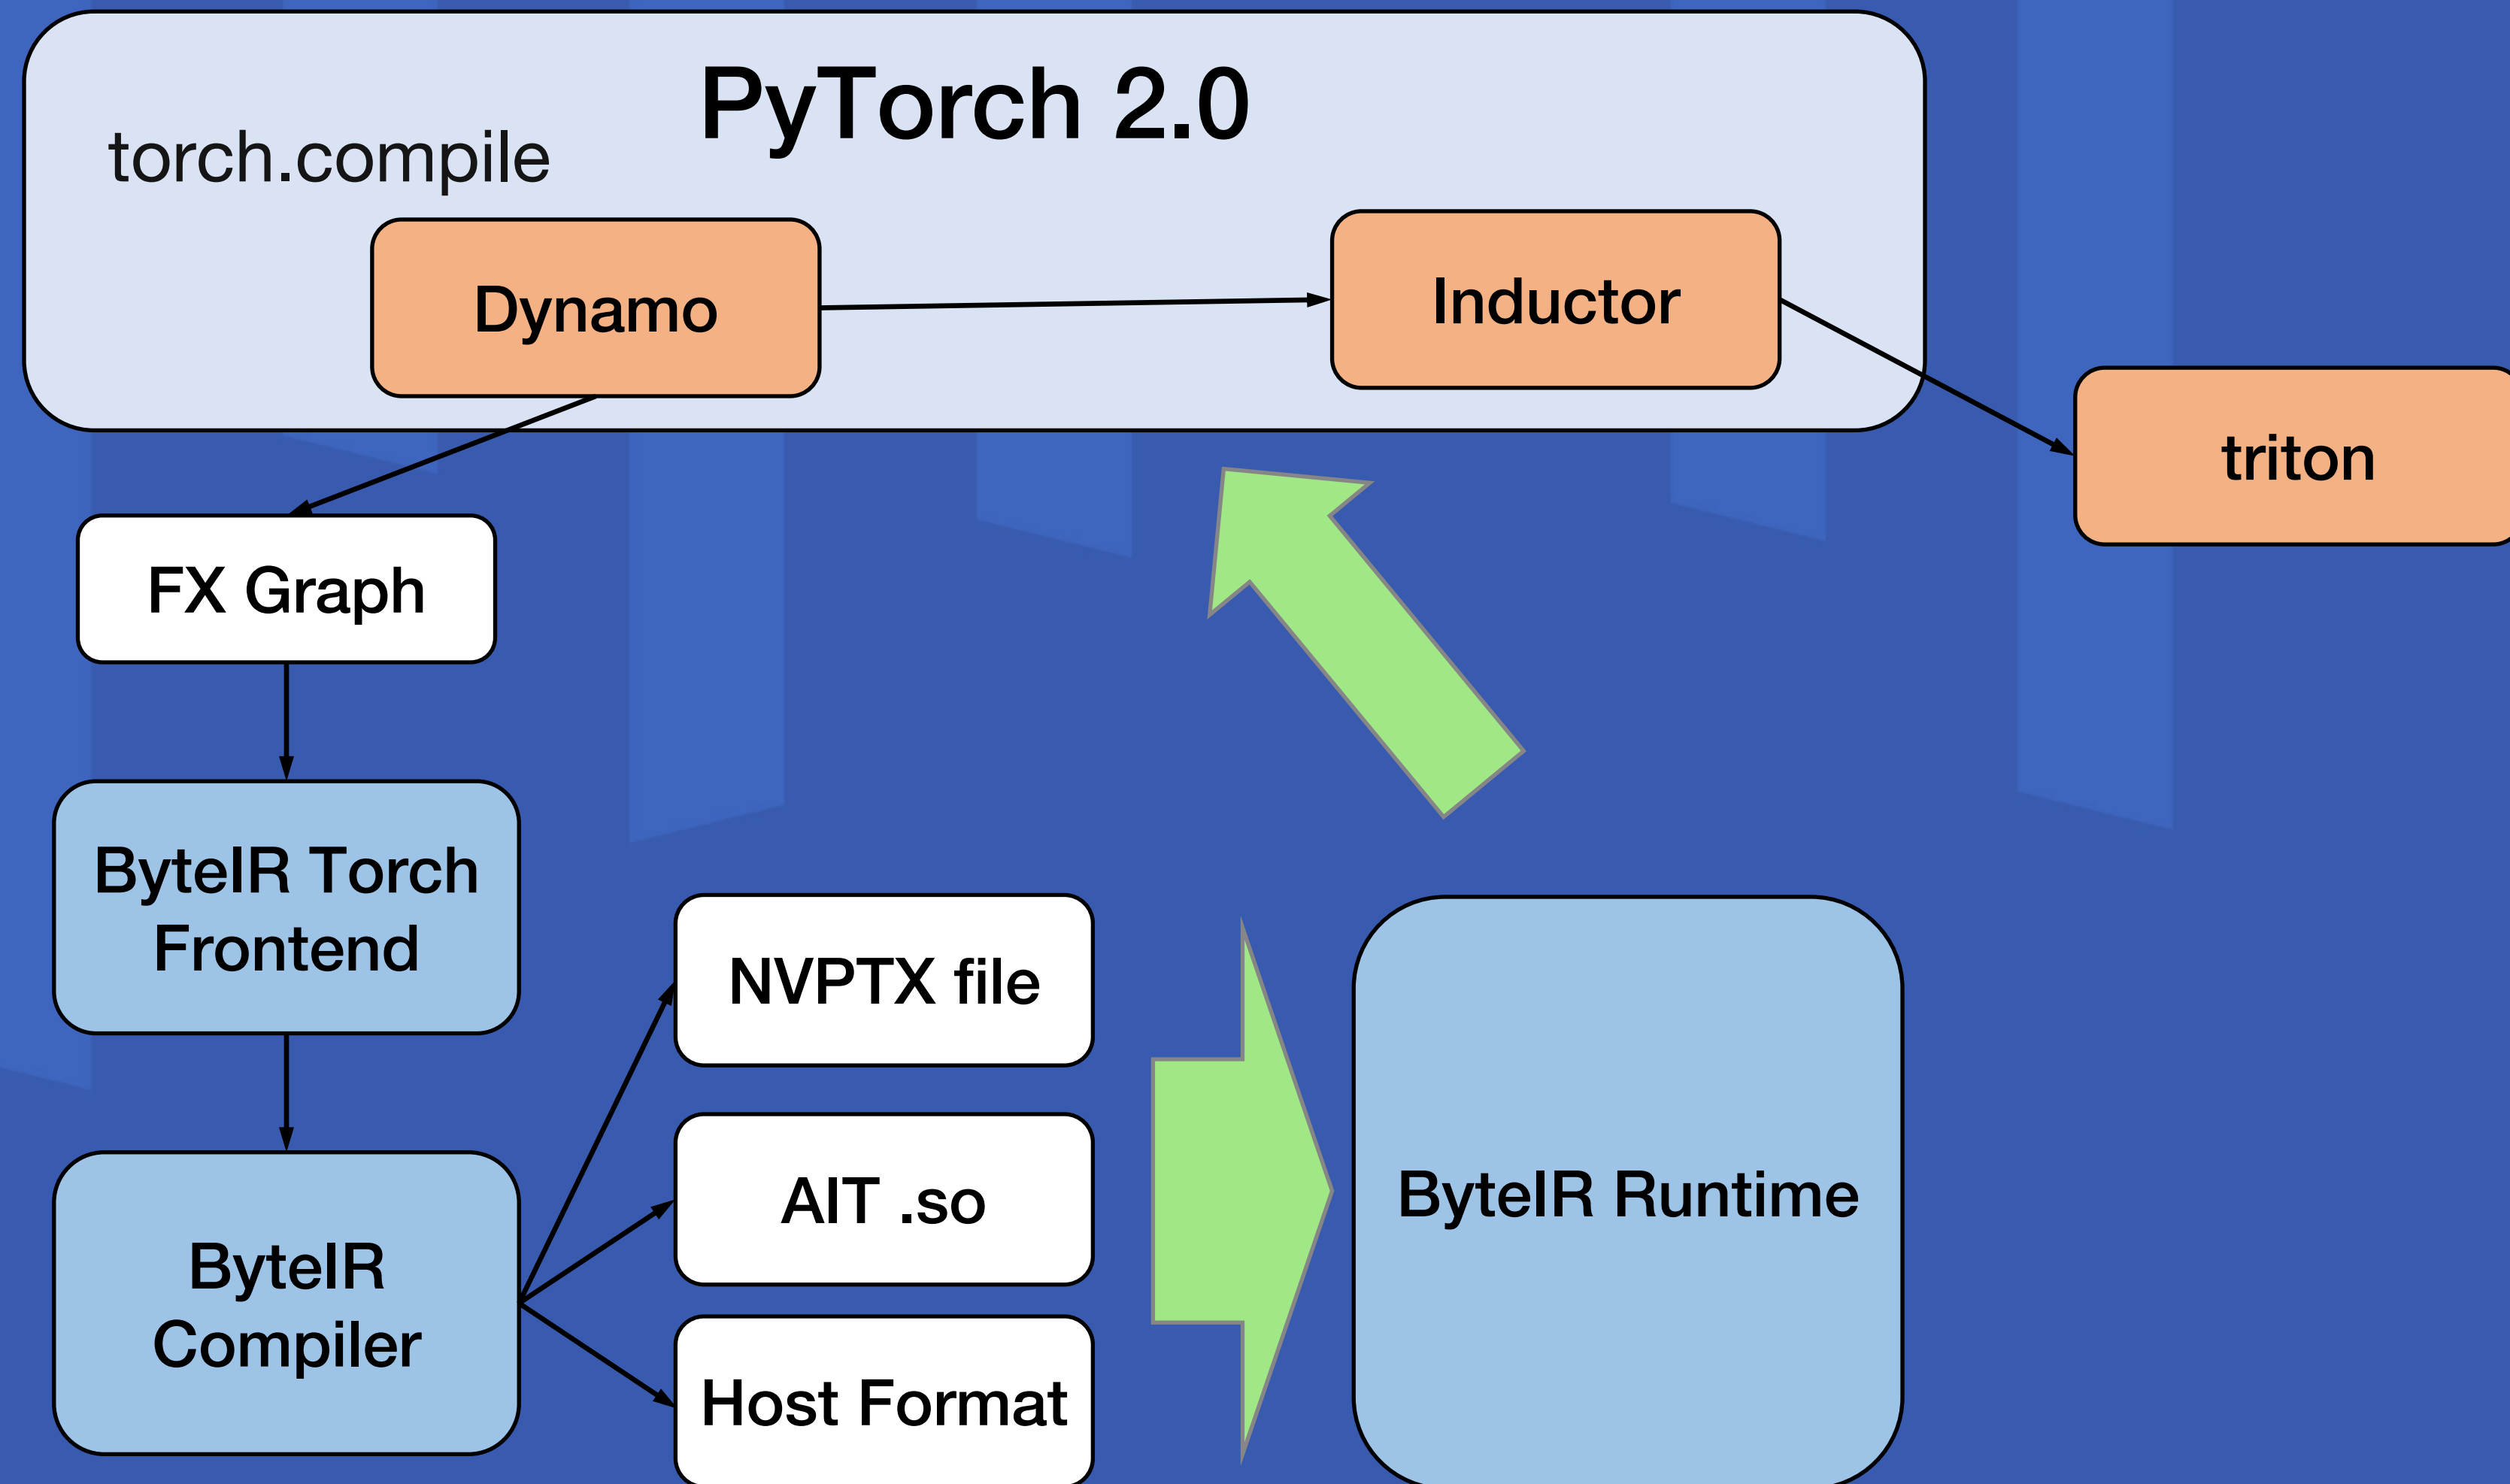
Memory Allocate/Free



1. What is BytelR and advantages of BytelR
2. Design and technical details
3. LLM training example and performance

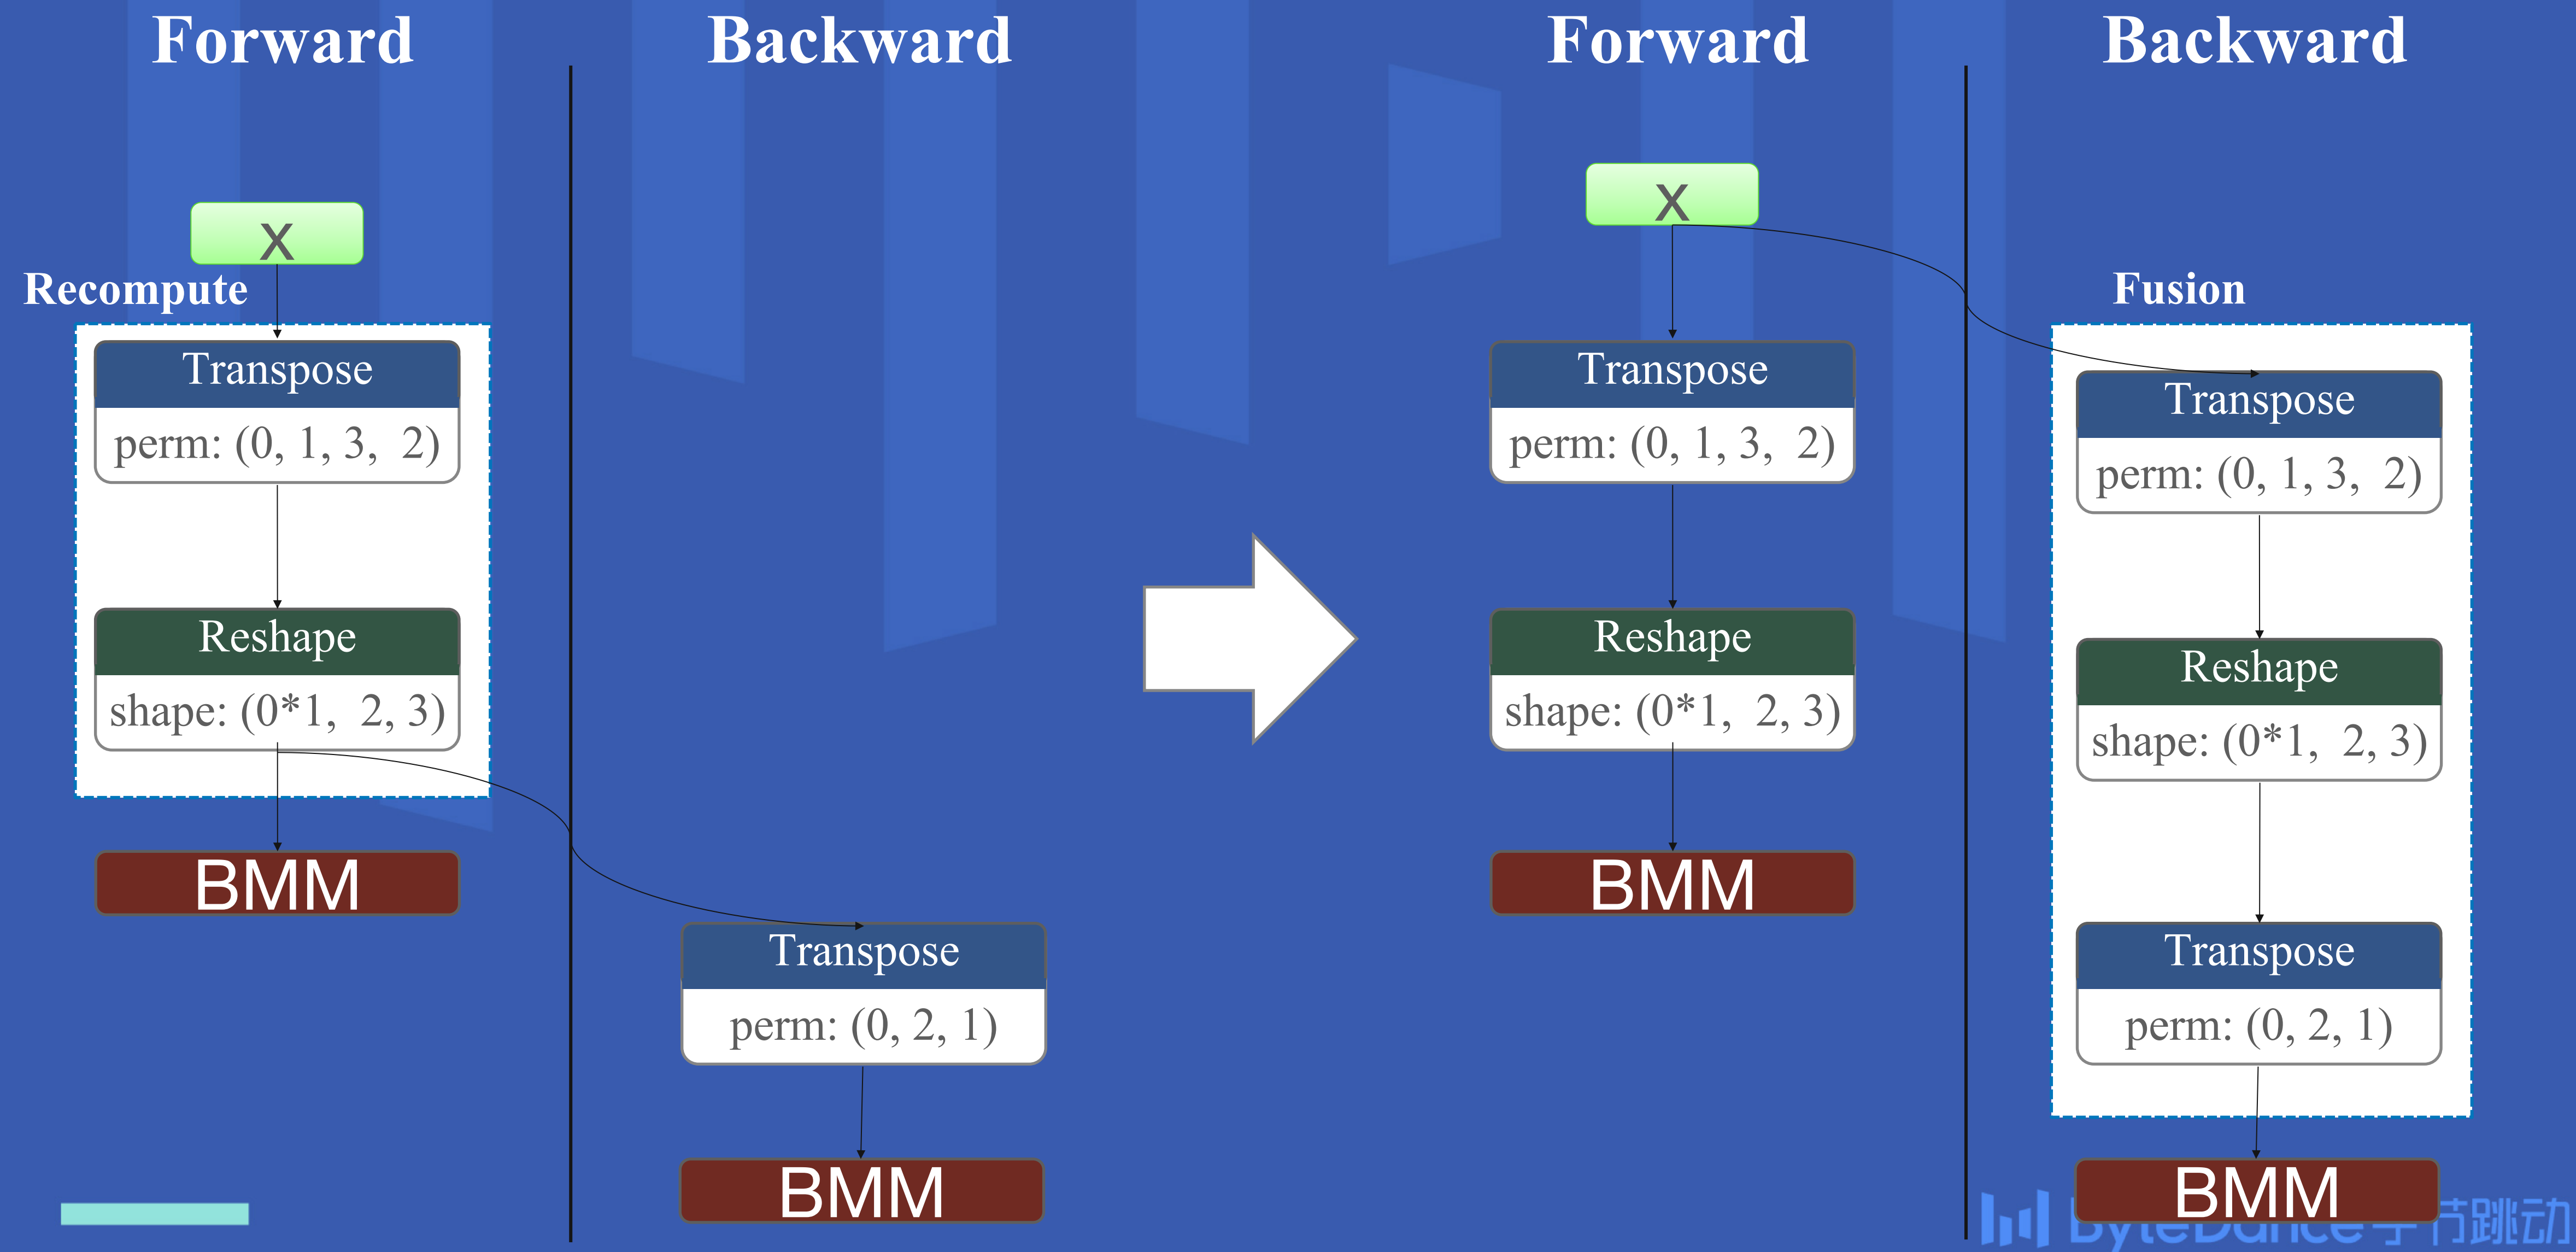


ByteIR LLM Training Compilation Pipeline



Demo Link: <https://github.com/bytedance/byteir/tree/main/frontends/torch-frontend/examples/demo>

Optimization in FW/BW Partition



One case of our partition strategy

ByteIR PyTorch Compile Example

```
from byteir import byteir_compile_fx

model = make_model(model_name)

# 1, compile with byteir
optimized_model = torch.compile(model, backend=byteir_compile_fx)

# 2, execution as usual
data = make_data(optimized_model, model_name, device)
model.zero_grad(set_to_none=True)
with torch.cuda.amp.autocast(enabled=True, dtype=torch.float16):
    # forward compile
    loss = compute_loss(optimized_model, data)
    # backward compile
    loss.backward()
```

ByteIR Runtime Example

```
import brt


session = brt.Session()
session.load(byre_model_path)
req = session.new_request_context(torch.cuda.current_stream())
inputs, outputs = [], []

# init input/output data
for offset in session.get_input_arg_offsets():
    inputs.append(torch.randn(session.get_static_shape(offset),
                              dtype=dtype, device="cuda"))
    req.bind_arg(offset, inputs[-1].data_ptr())
...

req.finish_io_binding()
req.run()
req.sync()
```


Performance

- Flash Attention 2
- Elementwise Fusion
- AITemplate
- Reduce Codegen



Conclusion & Future Work

We introduce ByteIR: a framework-to-hardware compiler solution

- Friendly to PyTorch and GPU/ASIC
- PyTorch 2.0 training/inference demo on LLM

Future Work:

- Distributed support
- TensorCore MMA Codegen

Website: <https://byteir.ai>

Github: <https://github.com/bytedance/byteir>



THANKS

<https://byteir.ai>

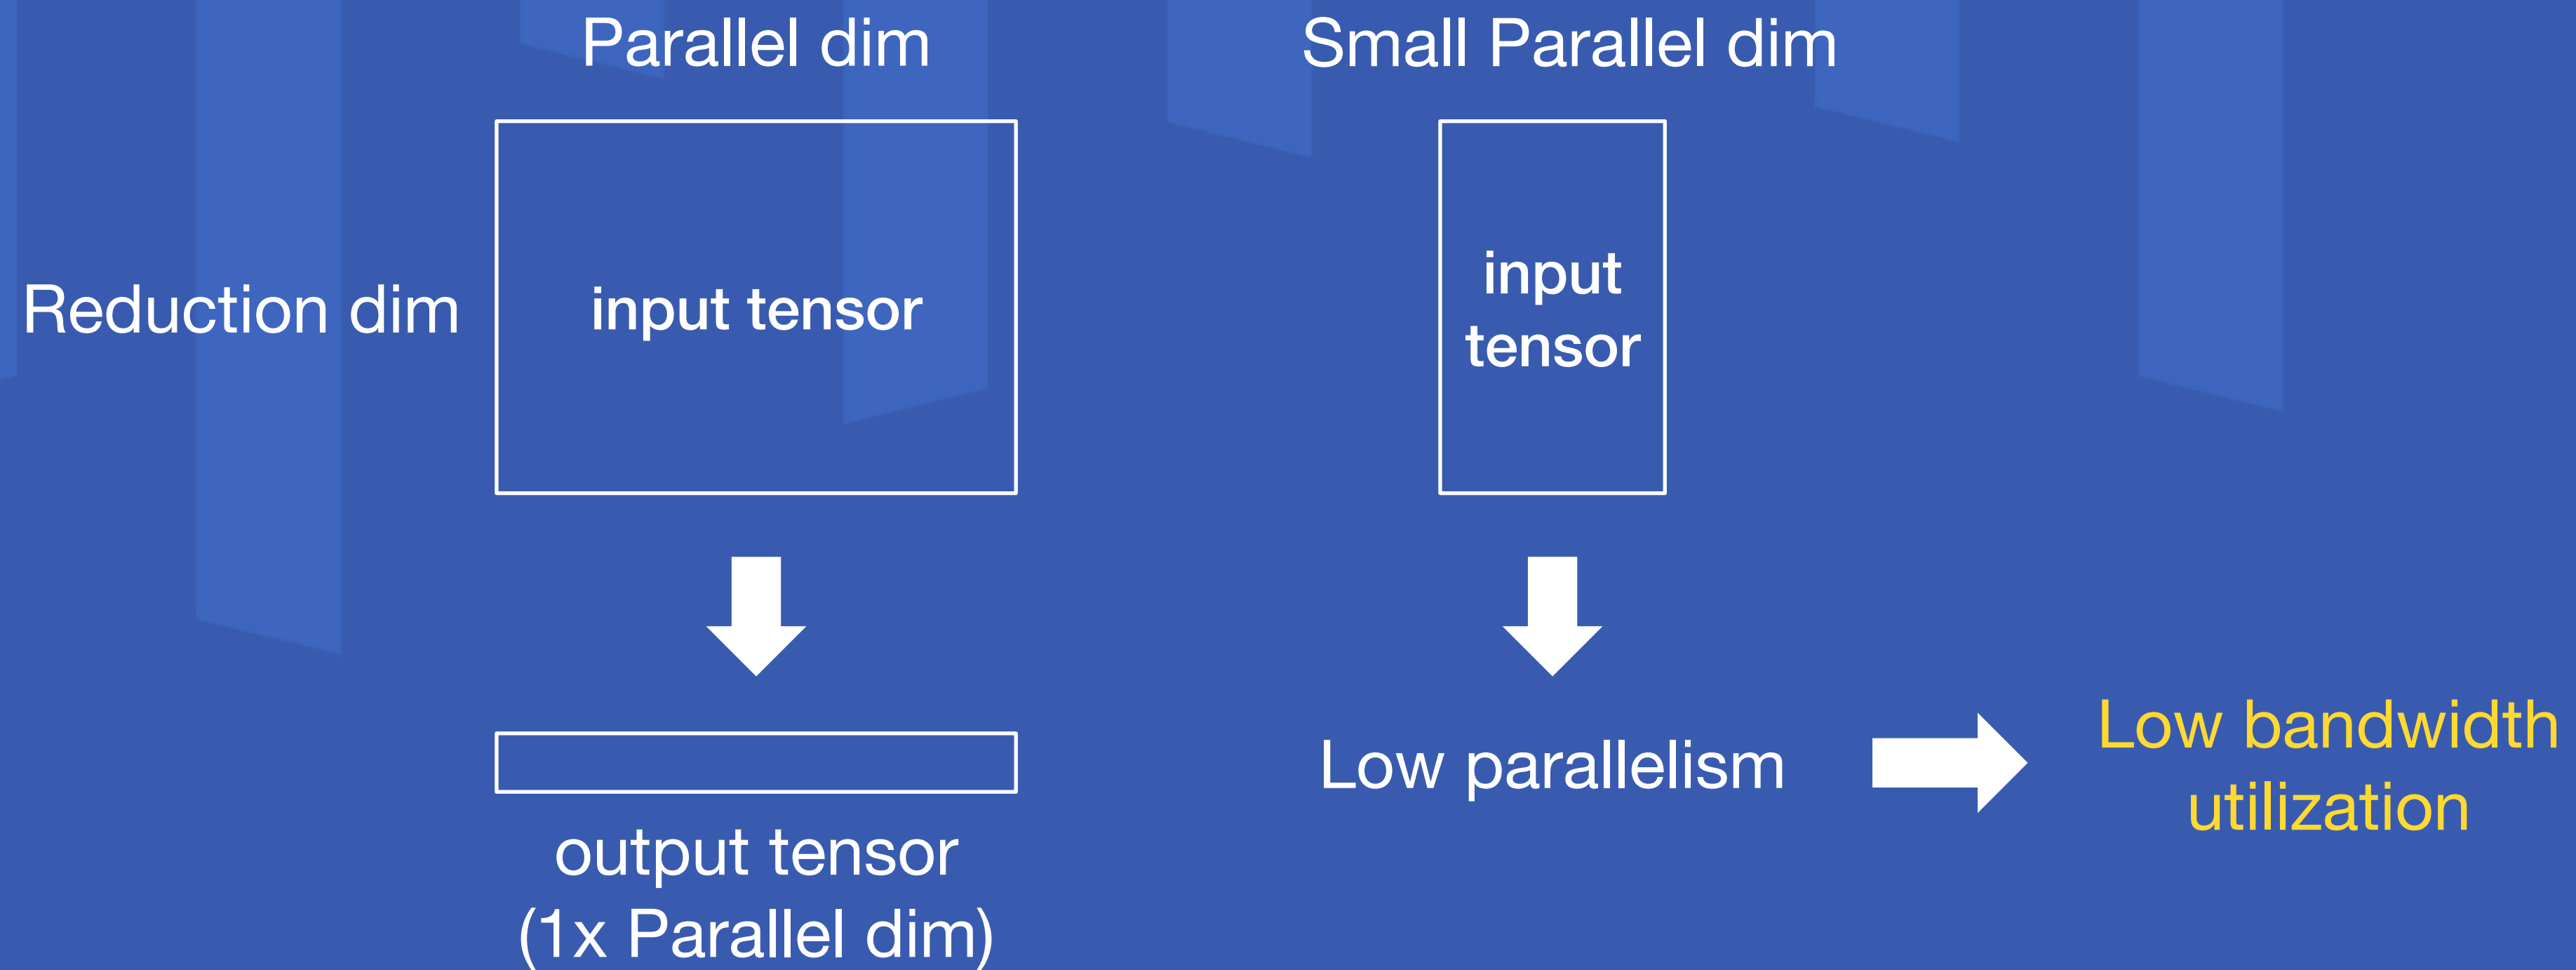
ByteIR Github: <https://github.com/bytedance/byteir>

Personal Github: <https://github.com/qingyunqu>



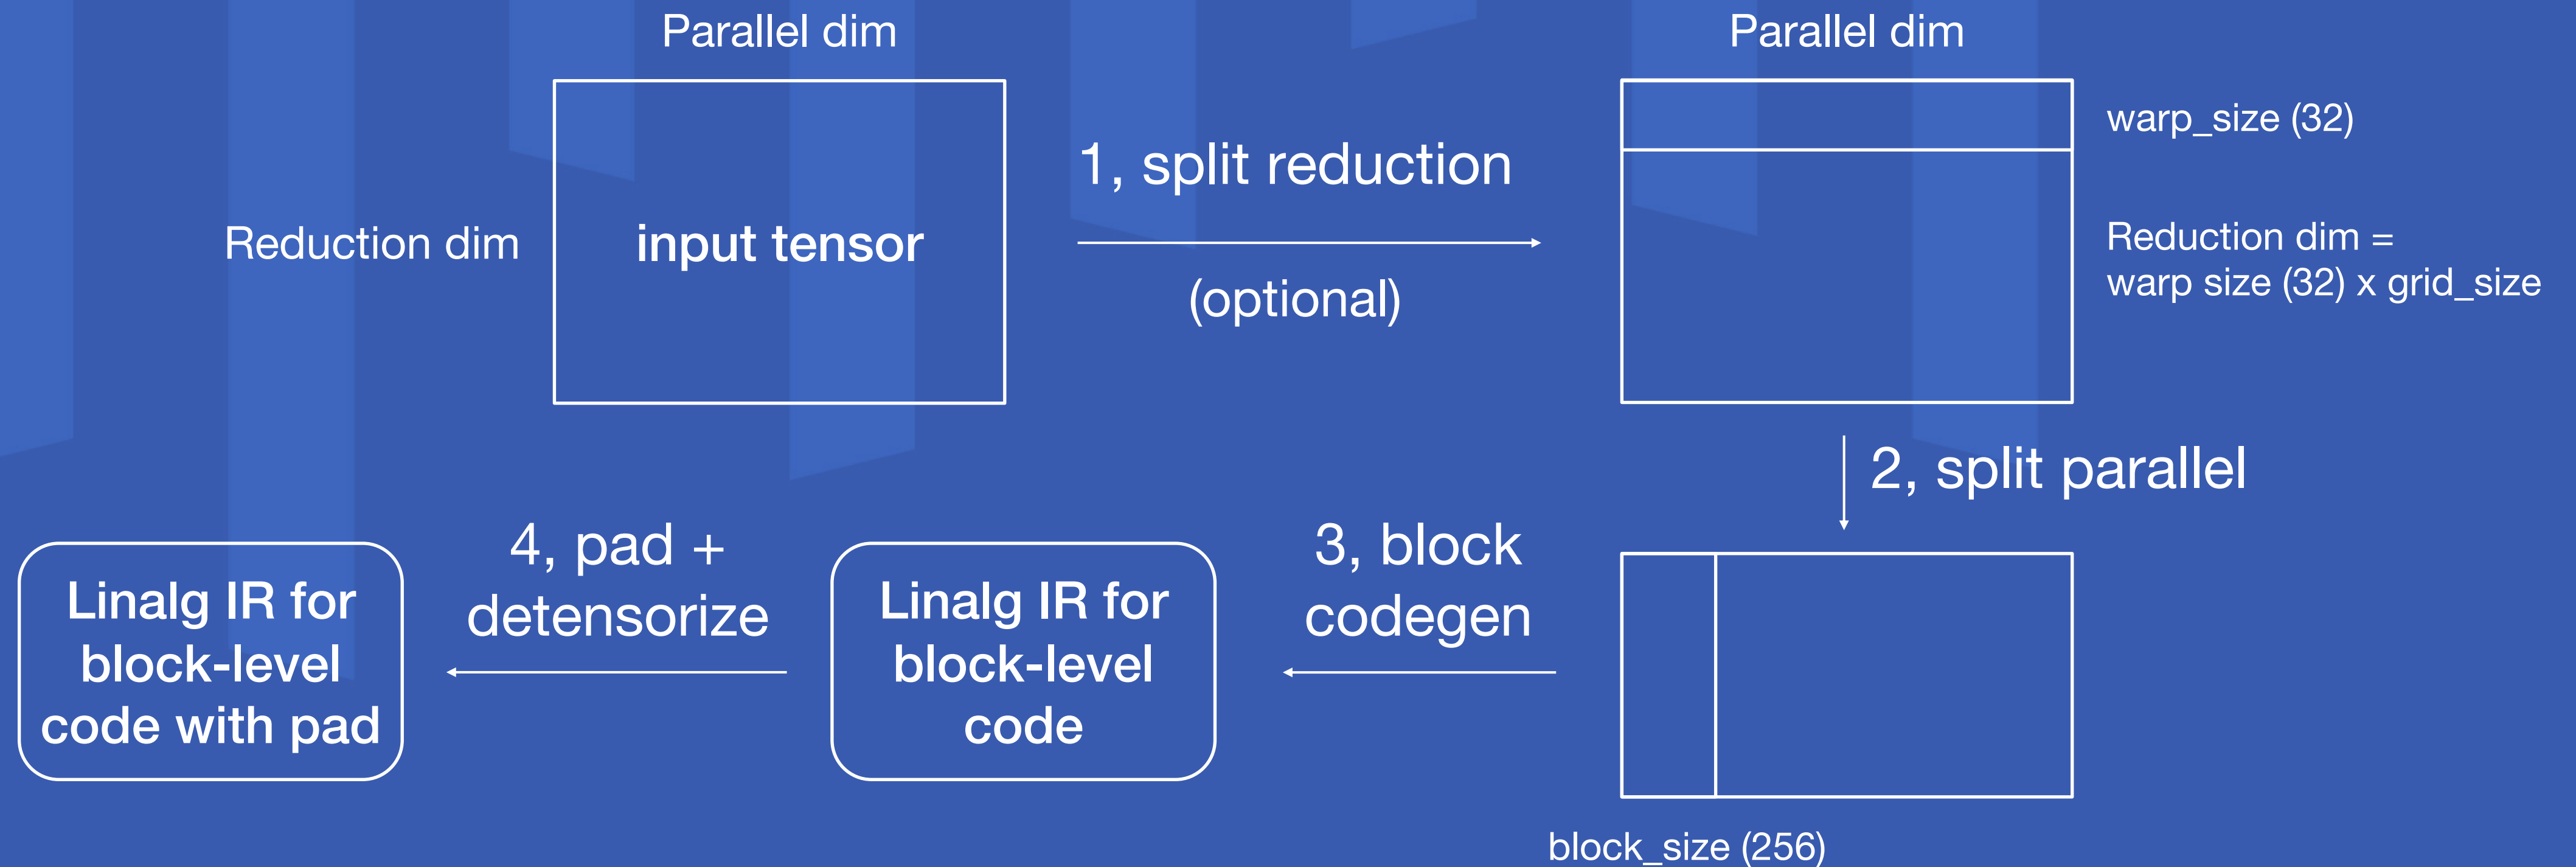
Appendix 1: Reduce Op Optimization (Fusion)

Optimization 1: Fusing reduce op with producer ops



Appendix 2: Reduce Op Optimization (Tiling)

Optimization 2: Parallelizing reduce dimension



We use utilize our LinalgExt transformations to achieve best tiling efficiency